



SAPIENZA
UNIVERSITÀ DI ROMA

Design and Implementation of a Replicated File Storage System Using Conflict-free Replicated Data Types

Department of Computer, Control and Management Engineering “Antonio Ruberti”
Master’s Degree in Engineering in Computer Science

Candidate: Giorgio Francescone (Matr. 1974461)
Thesis Advisor: Prof. Andrea Vitaletti

January 28, 2026



Introduction



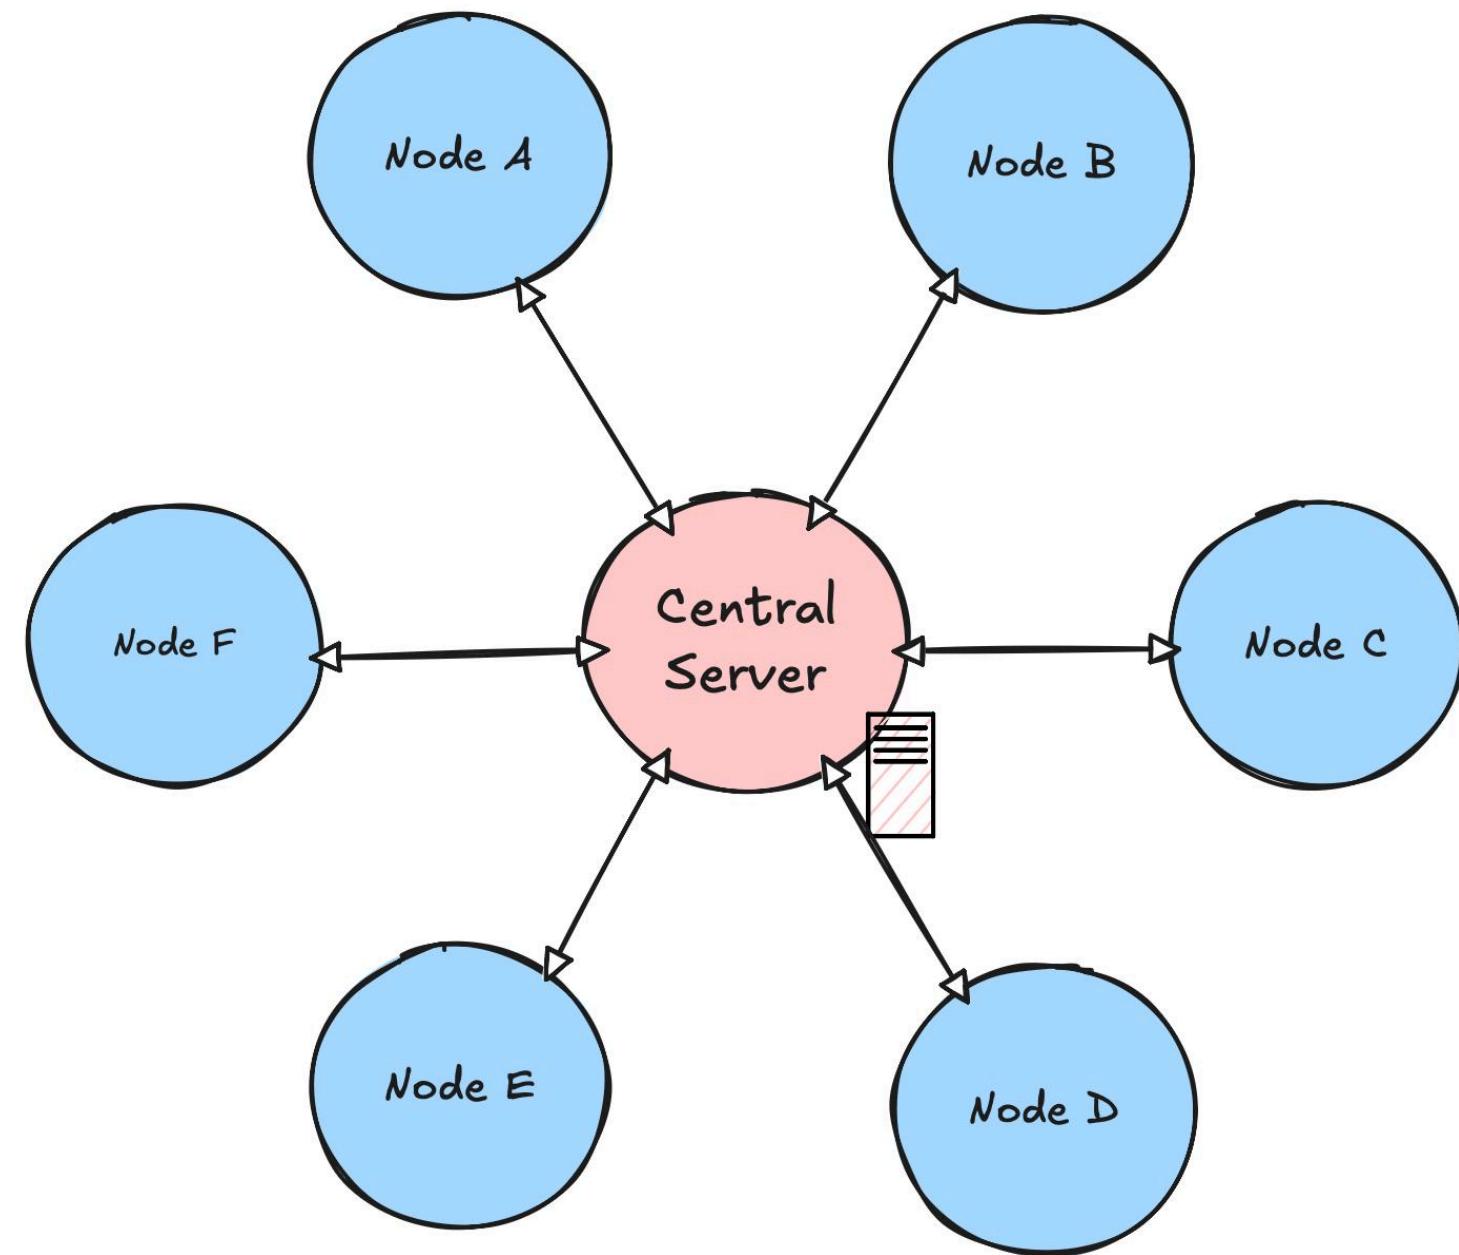
The shift towards distributed systems

- Modern systems distribute the workload onto **multiple machines** (as opposed to handling everything on a single, powerful machine)
- **Key benefits:** scalability, availability, efficiency

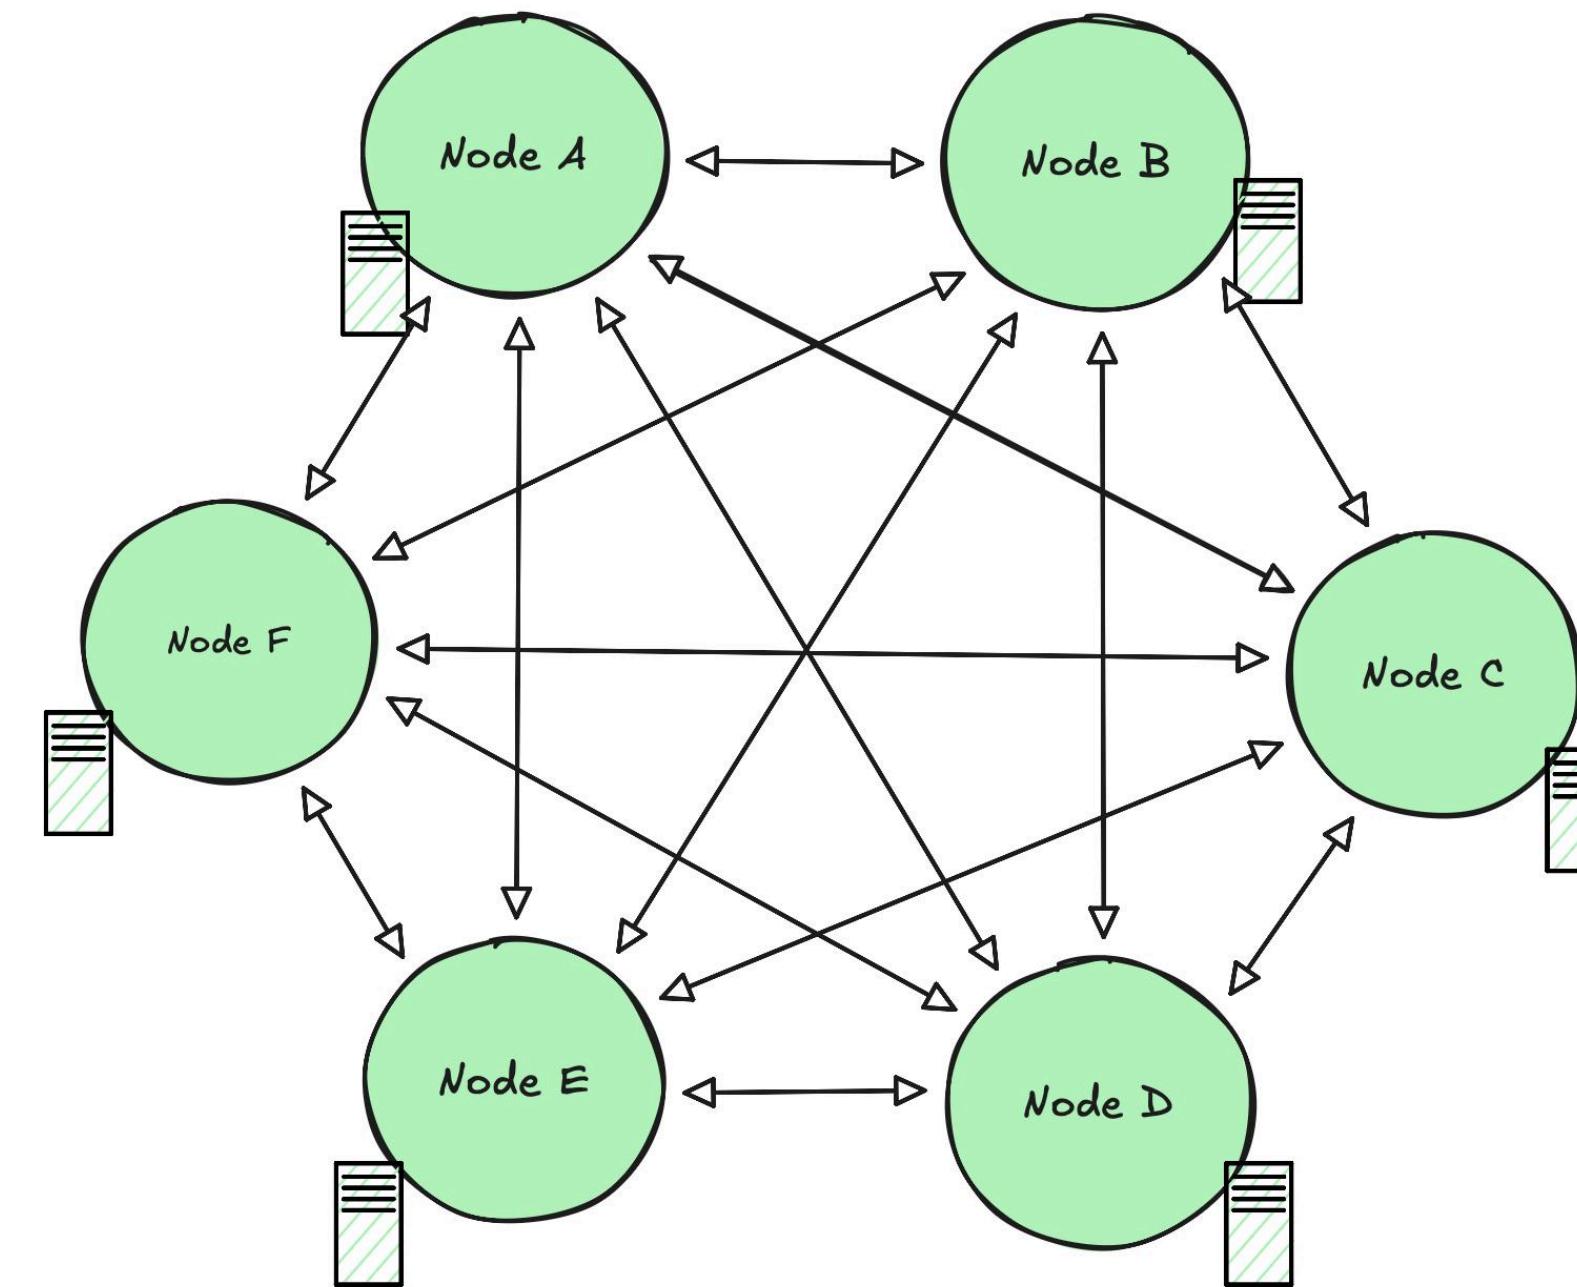
The challenges of collaborative applications

- Real-time, concurrent editing of shared data by multiple users (e.g., Google Docs)
 - How is everything kept **consistent**?
 - How to seamlessly handle *conflicts*?

Approaches to Replication

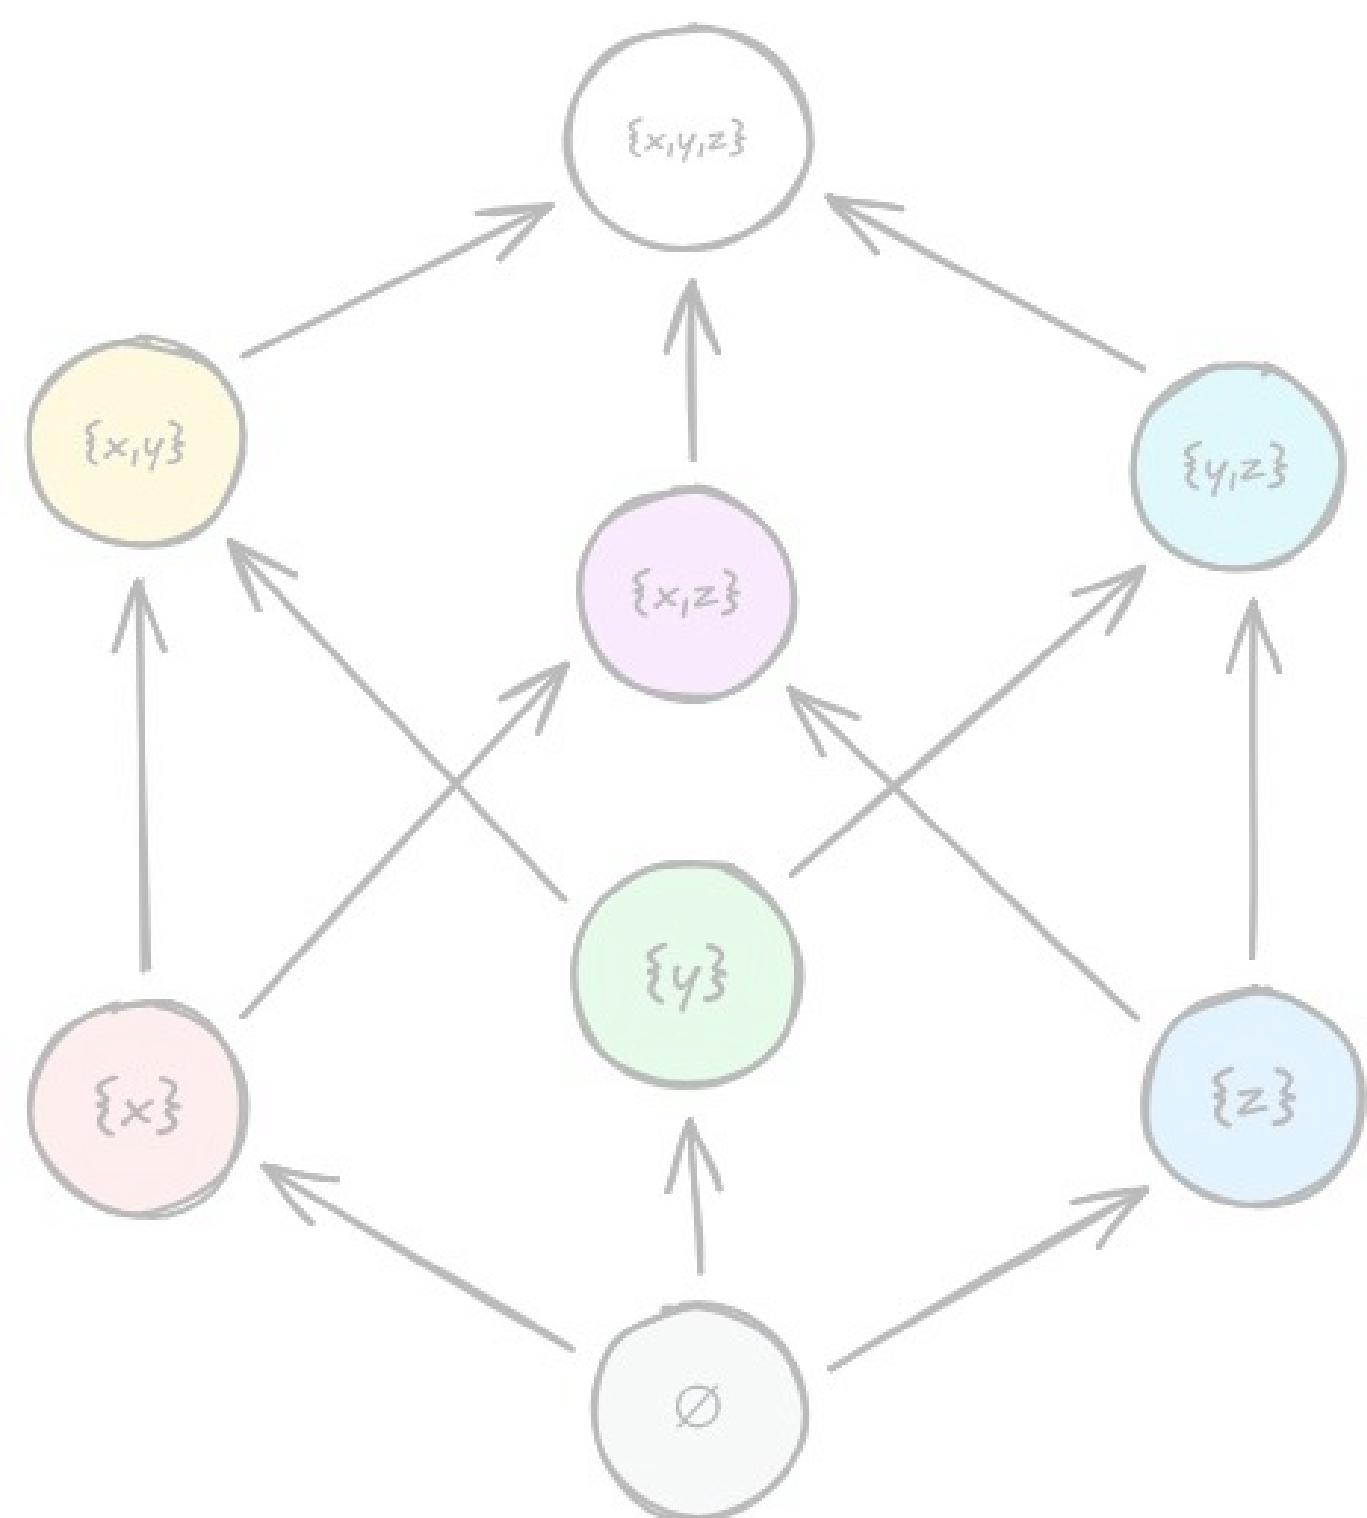


Centralized replication: a central server coordinates operations for the entire network and resolves conflicts



Decentralized replication: Each node individually handles (i.e. merges) operations received from other nodes — no need for coordination

Foundational Concepts on CRDTs



Strong Consistency

All nodes in the system see the same data at the same time

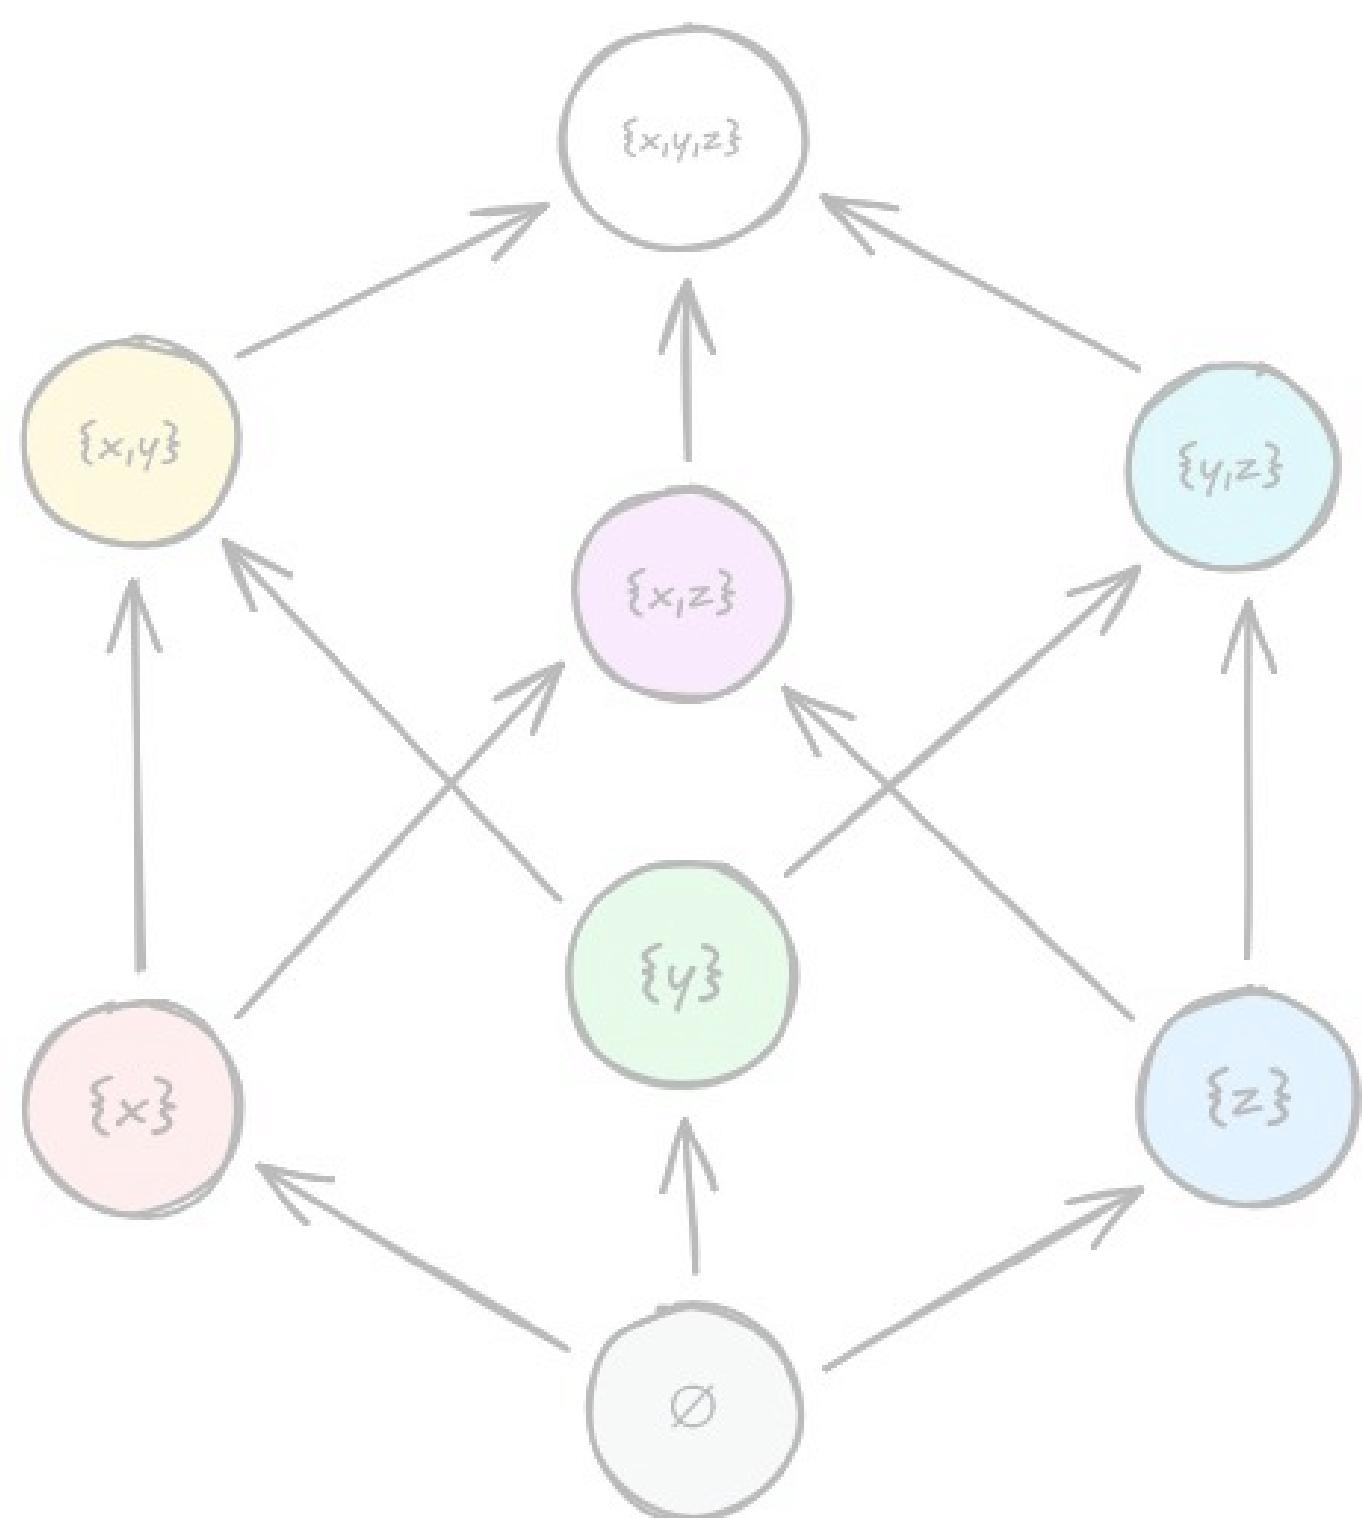
Eventual Consistency (EC)

Once updates are no longer being issued, all nodes in the system eventually converge to the same version of data

Strong Eventual Consistency (SEC)

Extension of EC which ensures *Strong Convergence* (nodes that receive the same updates converge to the same state)

Foundational Concepts on CRDTs



Strong Consistency

All nodes in the system see the same data at the same time

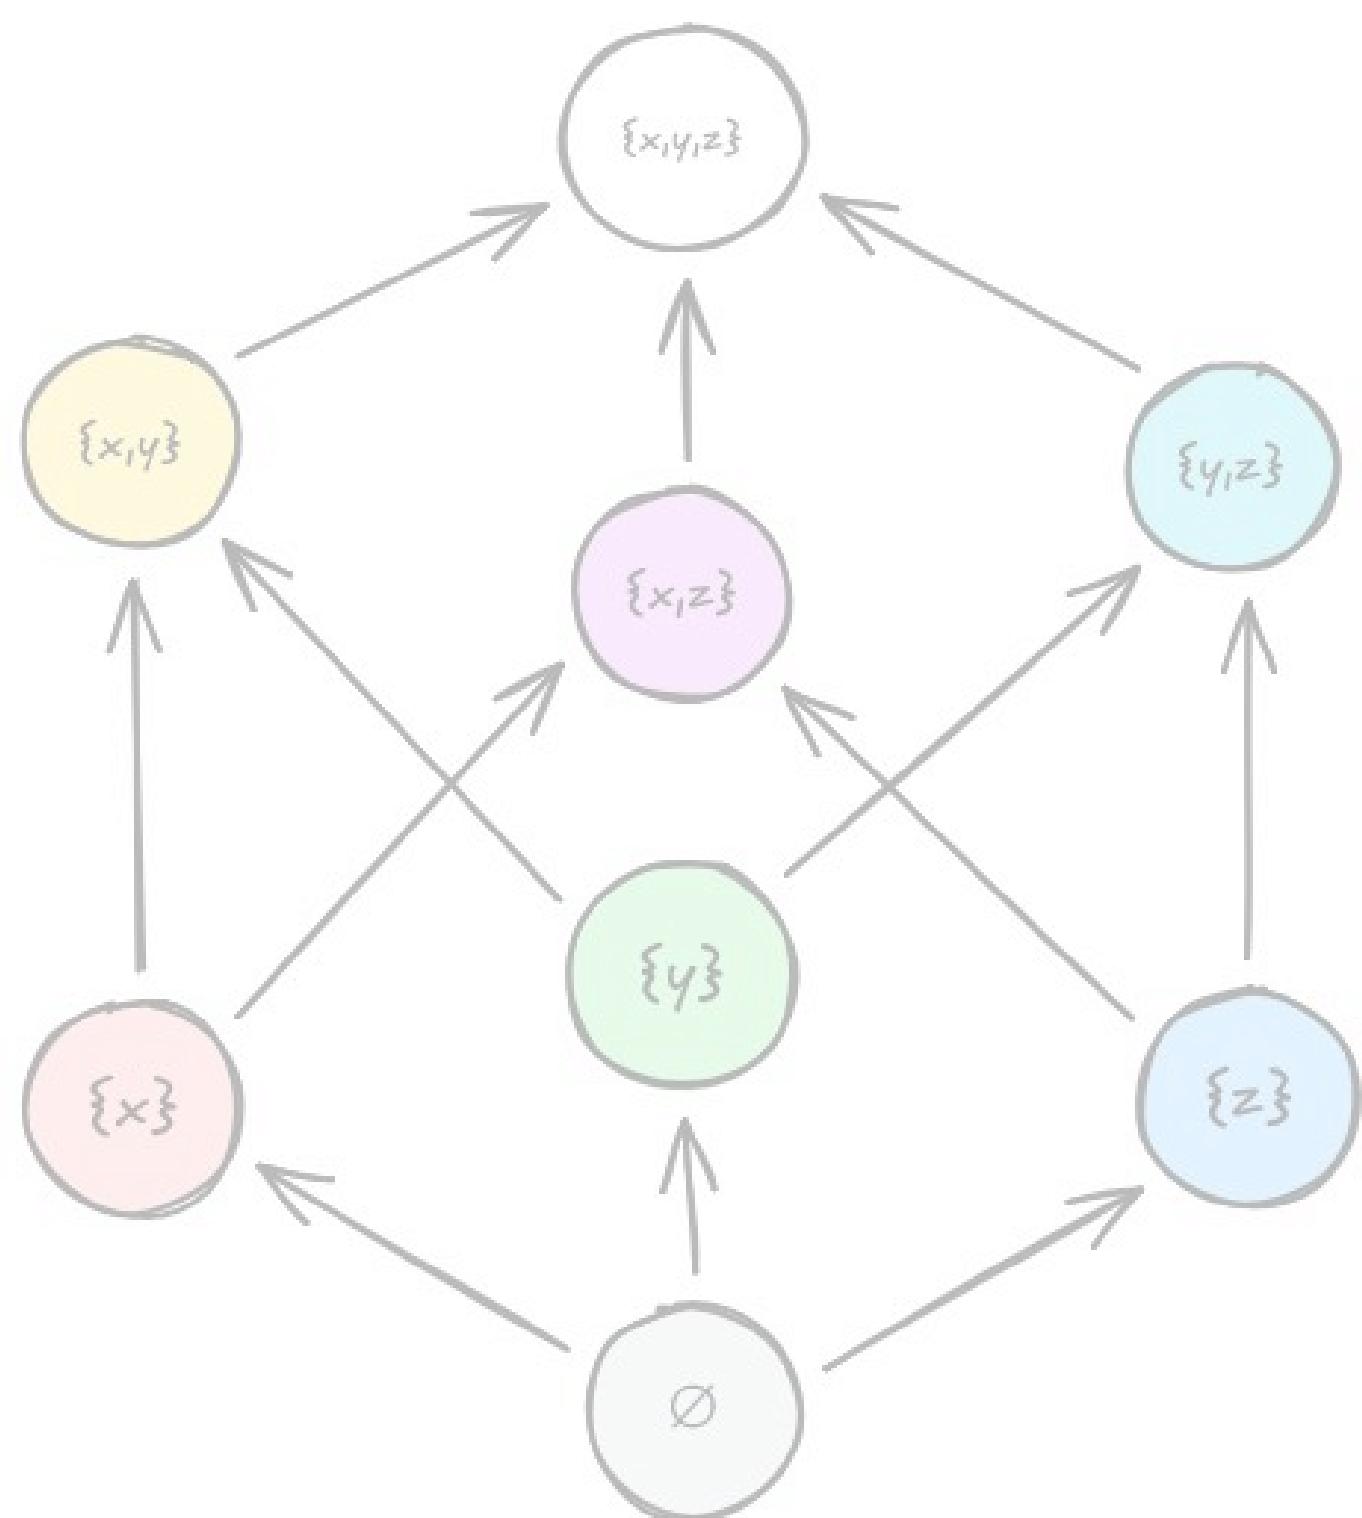
Eventual Consistency (EC)

Once updates are no longer being issued, all nodes in the system *eventually* converge to the same version of data

Strong Eventual Consistency (SEC)

Extension of EC which ensures *Strong Convergence* (nodes that receive the same updates converge to the same state)

Foundational Concepts on CRDTs



Strong Consistency

All nodes in the system see the same data at the same time

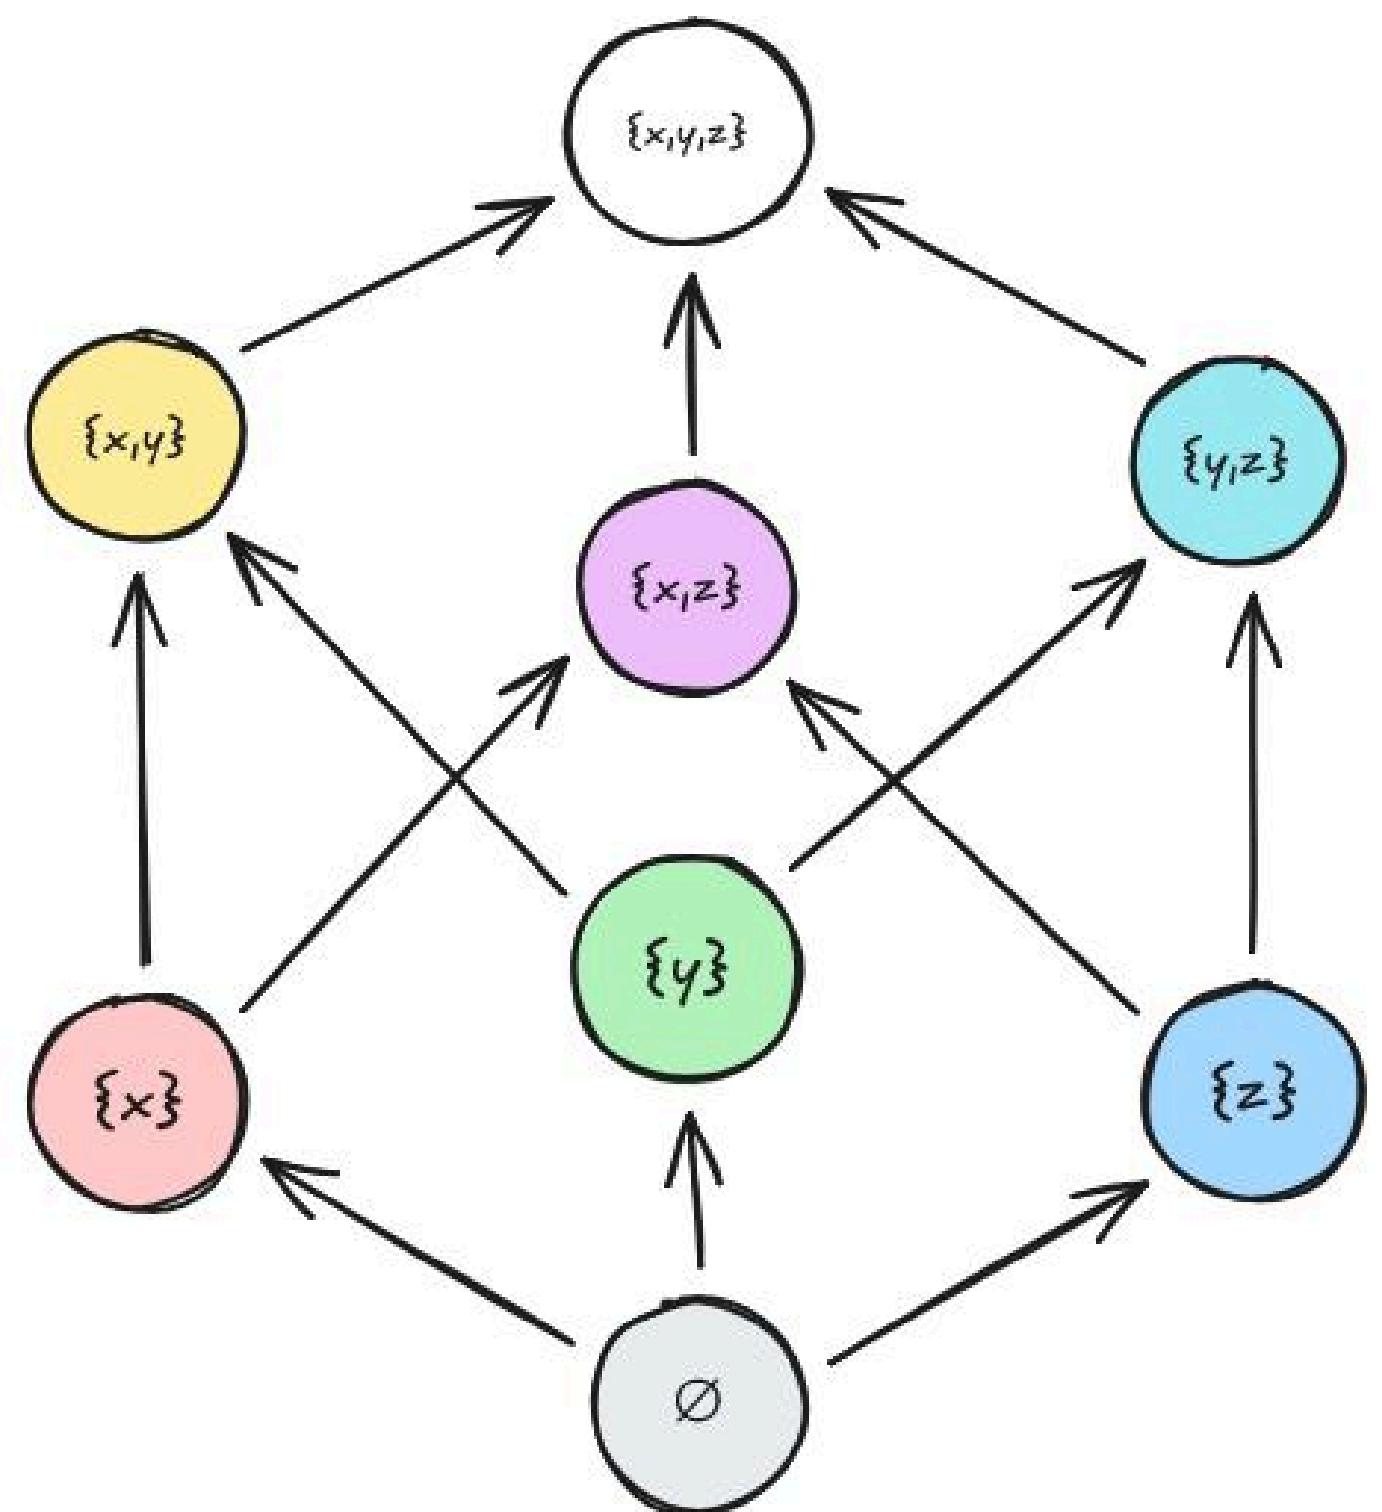
Eventual Consistency (EC)

Once updates are no longer being issued, all nodes in the system eventually converge to the same version of data

Strong Eventual Consistency (SEC)

Extension of EC which ensures *Strong Convergence* (nodes that receive the same updates converge to the same state)

Foundational Concepts on CRDTs

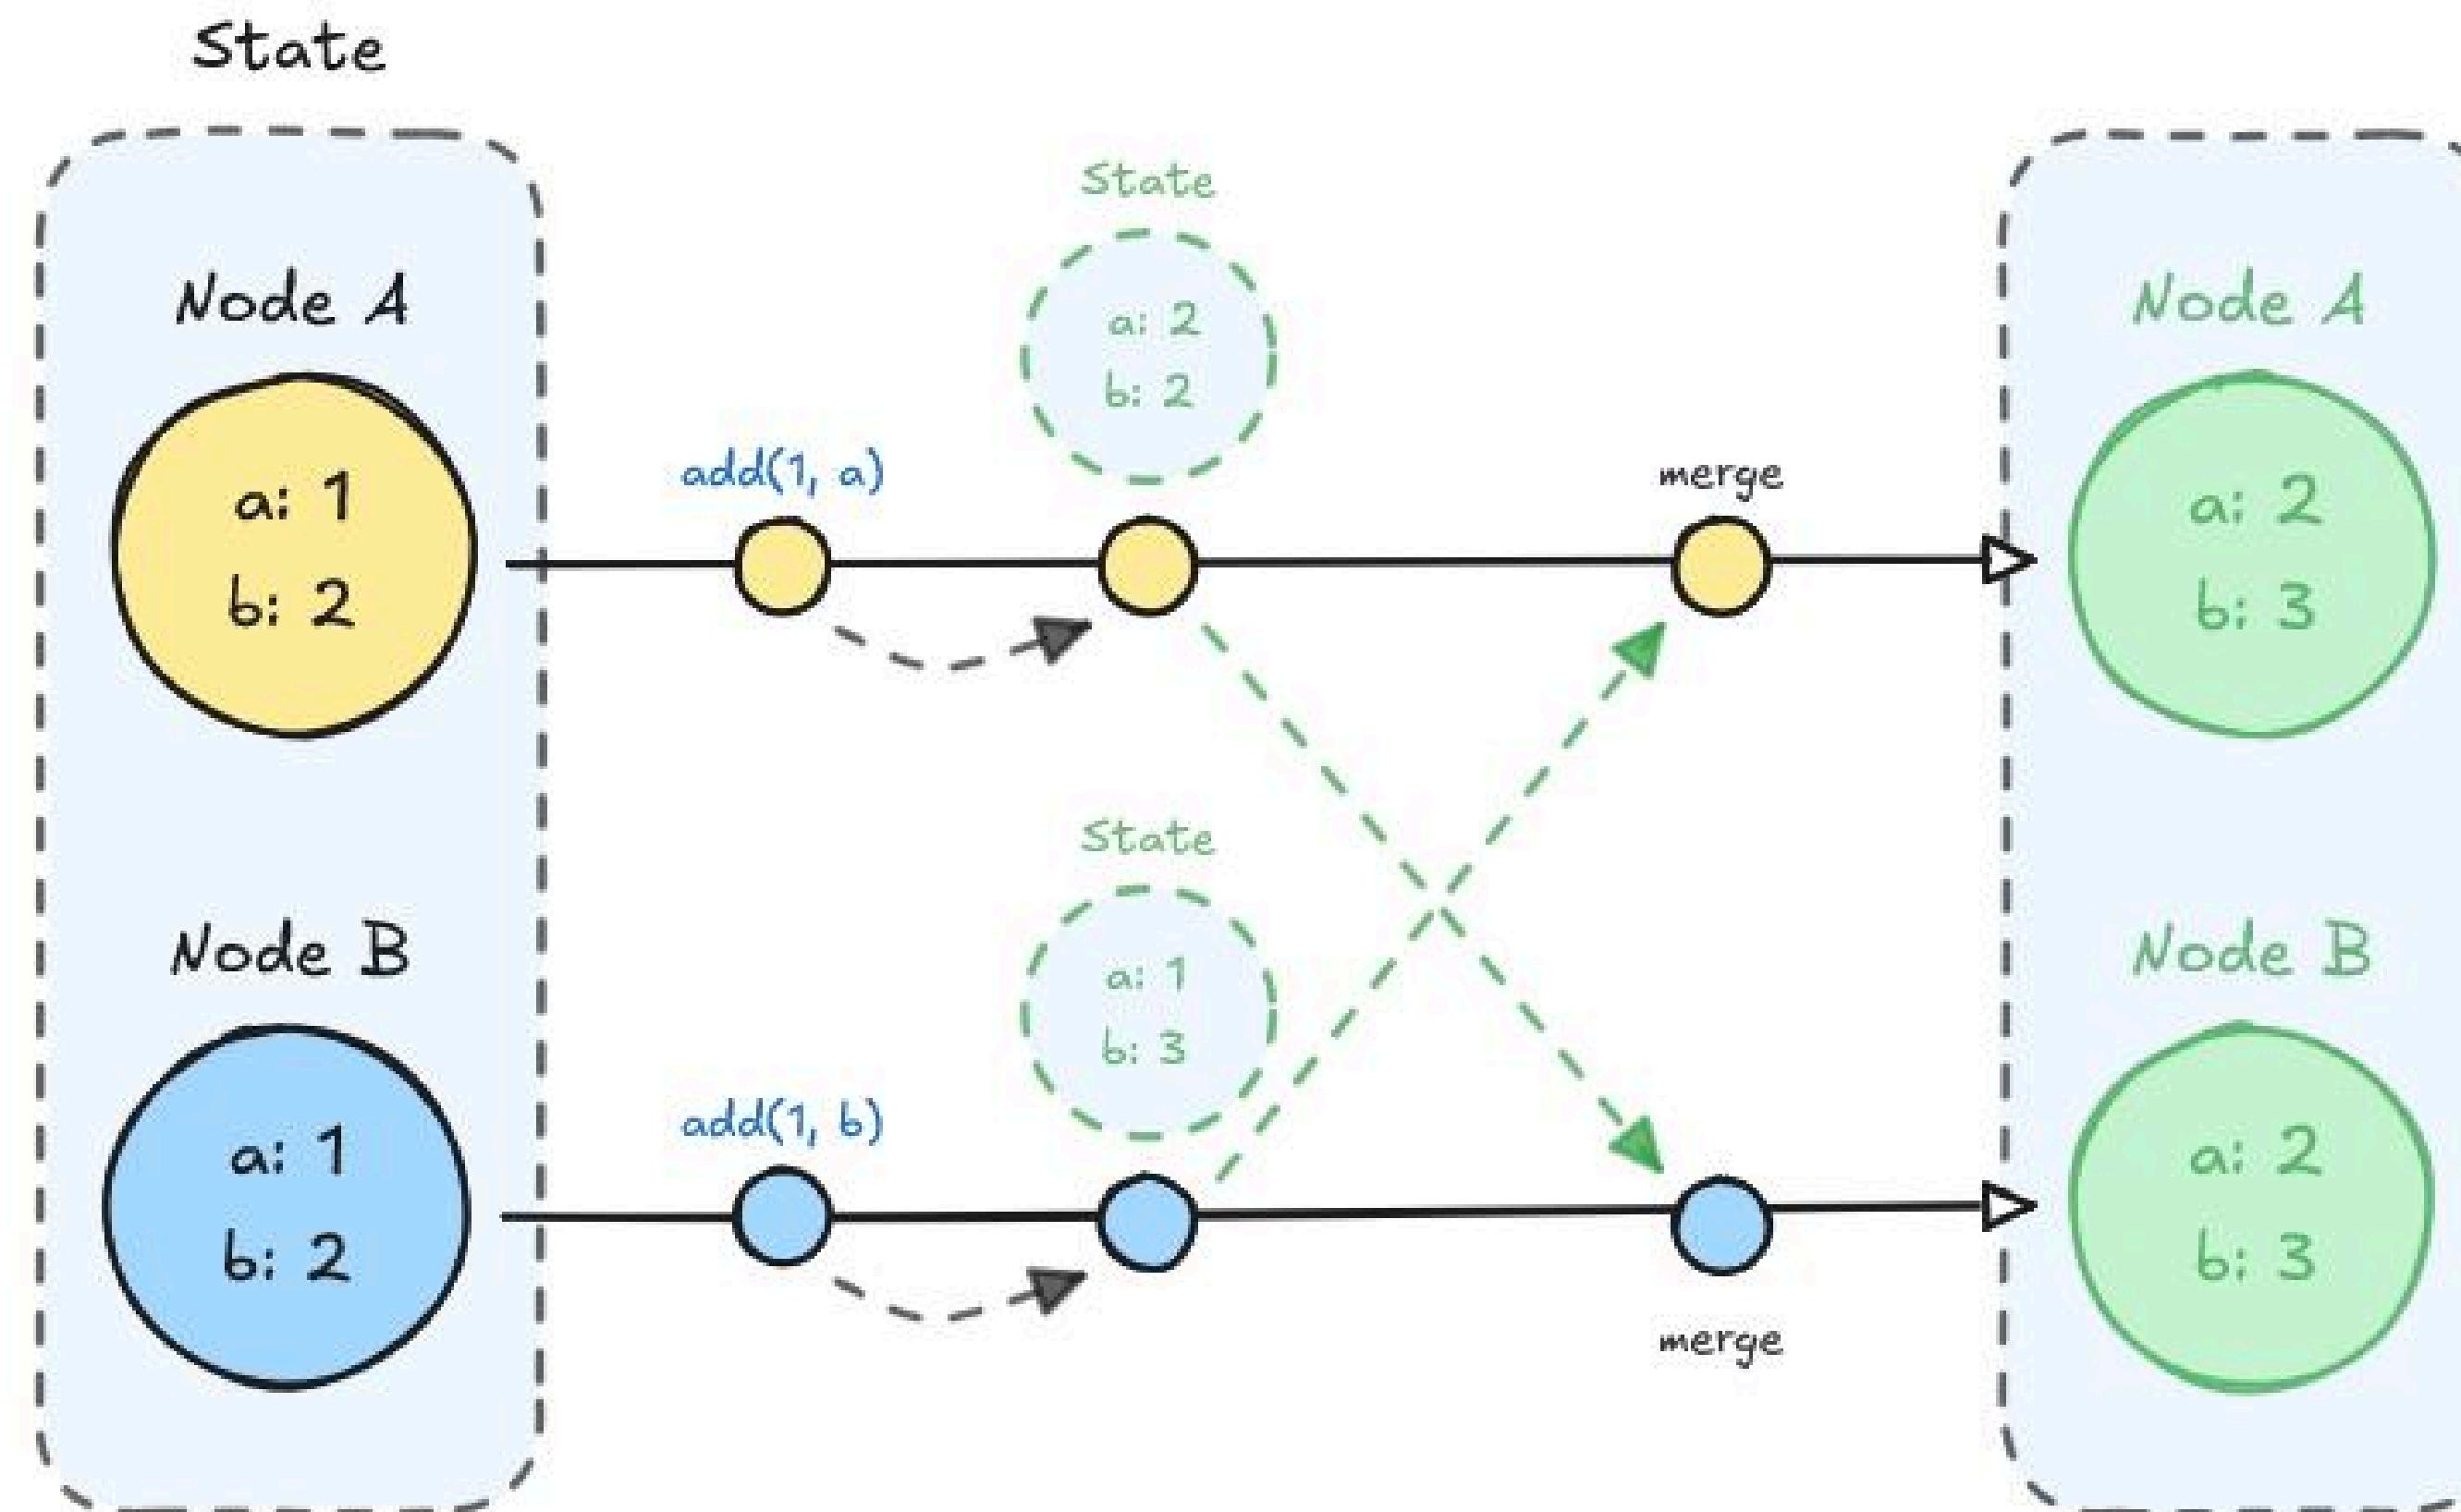


Conflict-free Replicated Data Types (CRDTs) are structures that enable *Strong Eventual Consistency* in a principled manner, allowing replicas in distributed systems to be updated independently and concurrently [1].

Replicas that have received the same state of updates will *eventually converge* to the same state, without requiring coordination or consensus protocols.

[1] M. Shapiro, N. Preguiça, C. Baquero, and M. Zawirski, “Conflict-free replicated data types”, In *Proceedings of the 13th International Conference on Stabilization, Safety, and Security of Distributed Systems*, 2011, pp. 386–400. doi: 10.1007/978-3-642-24550-3_29.

Types of CRDT



State-based CRDTs

Broadcast the entire state to other nodes.

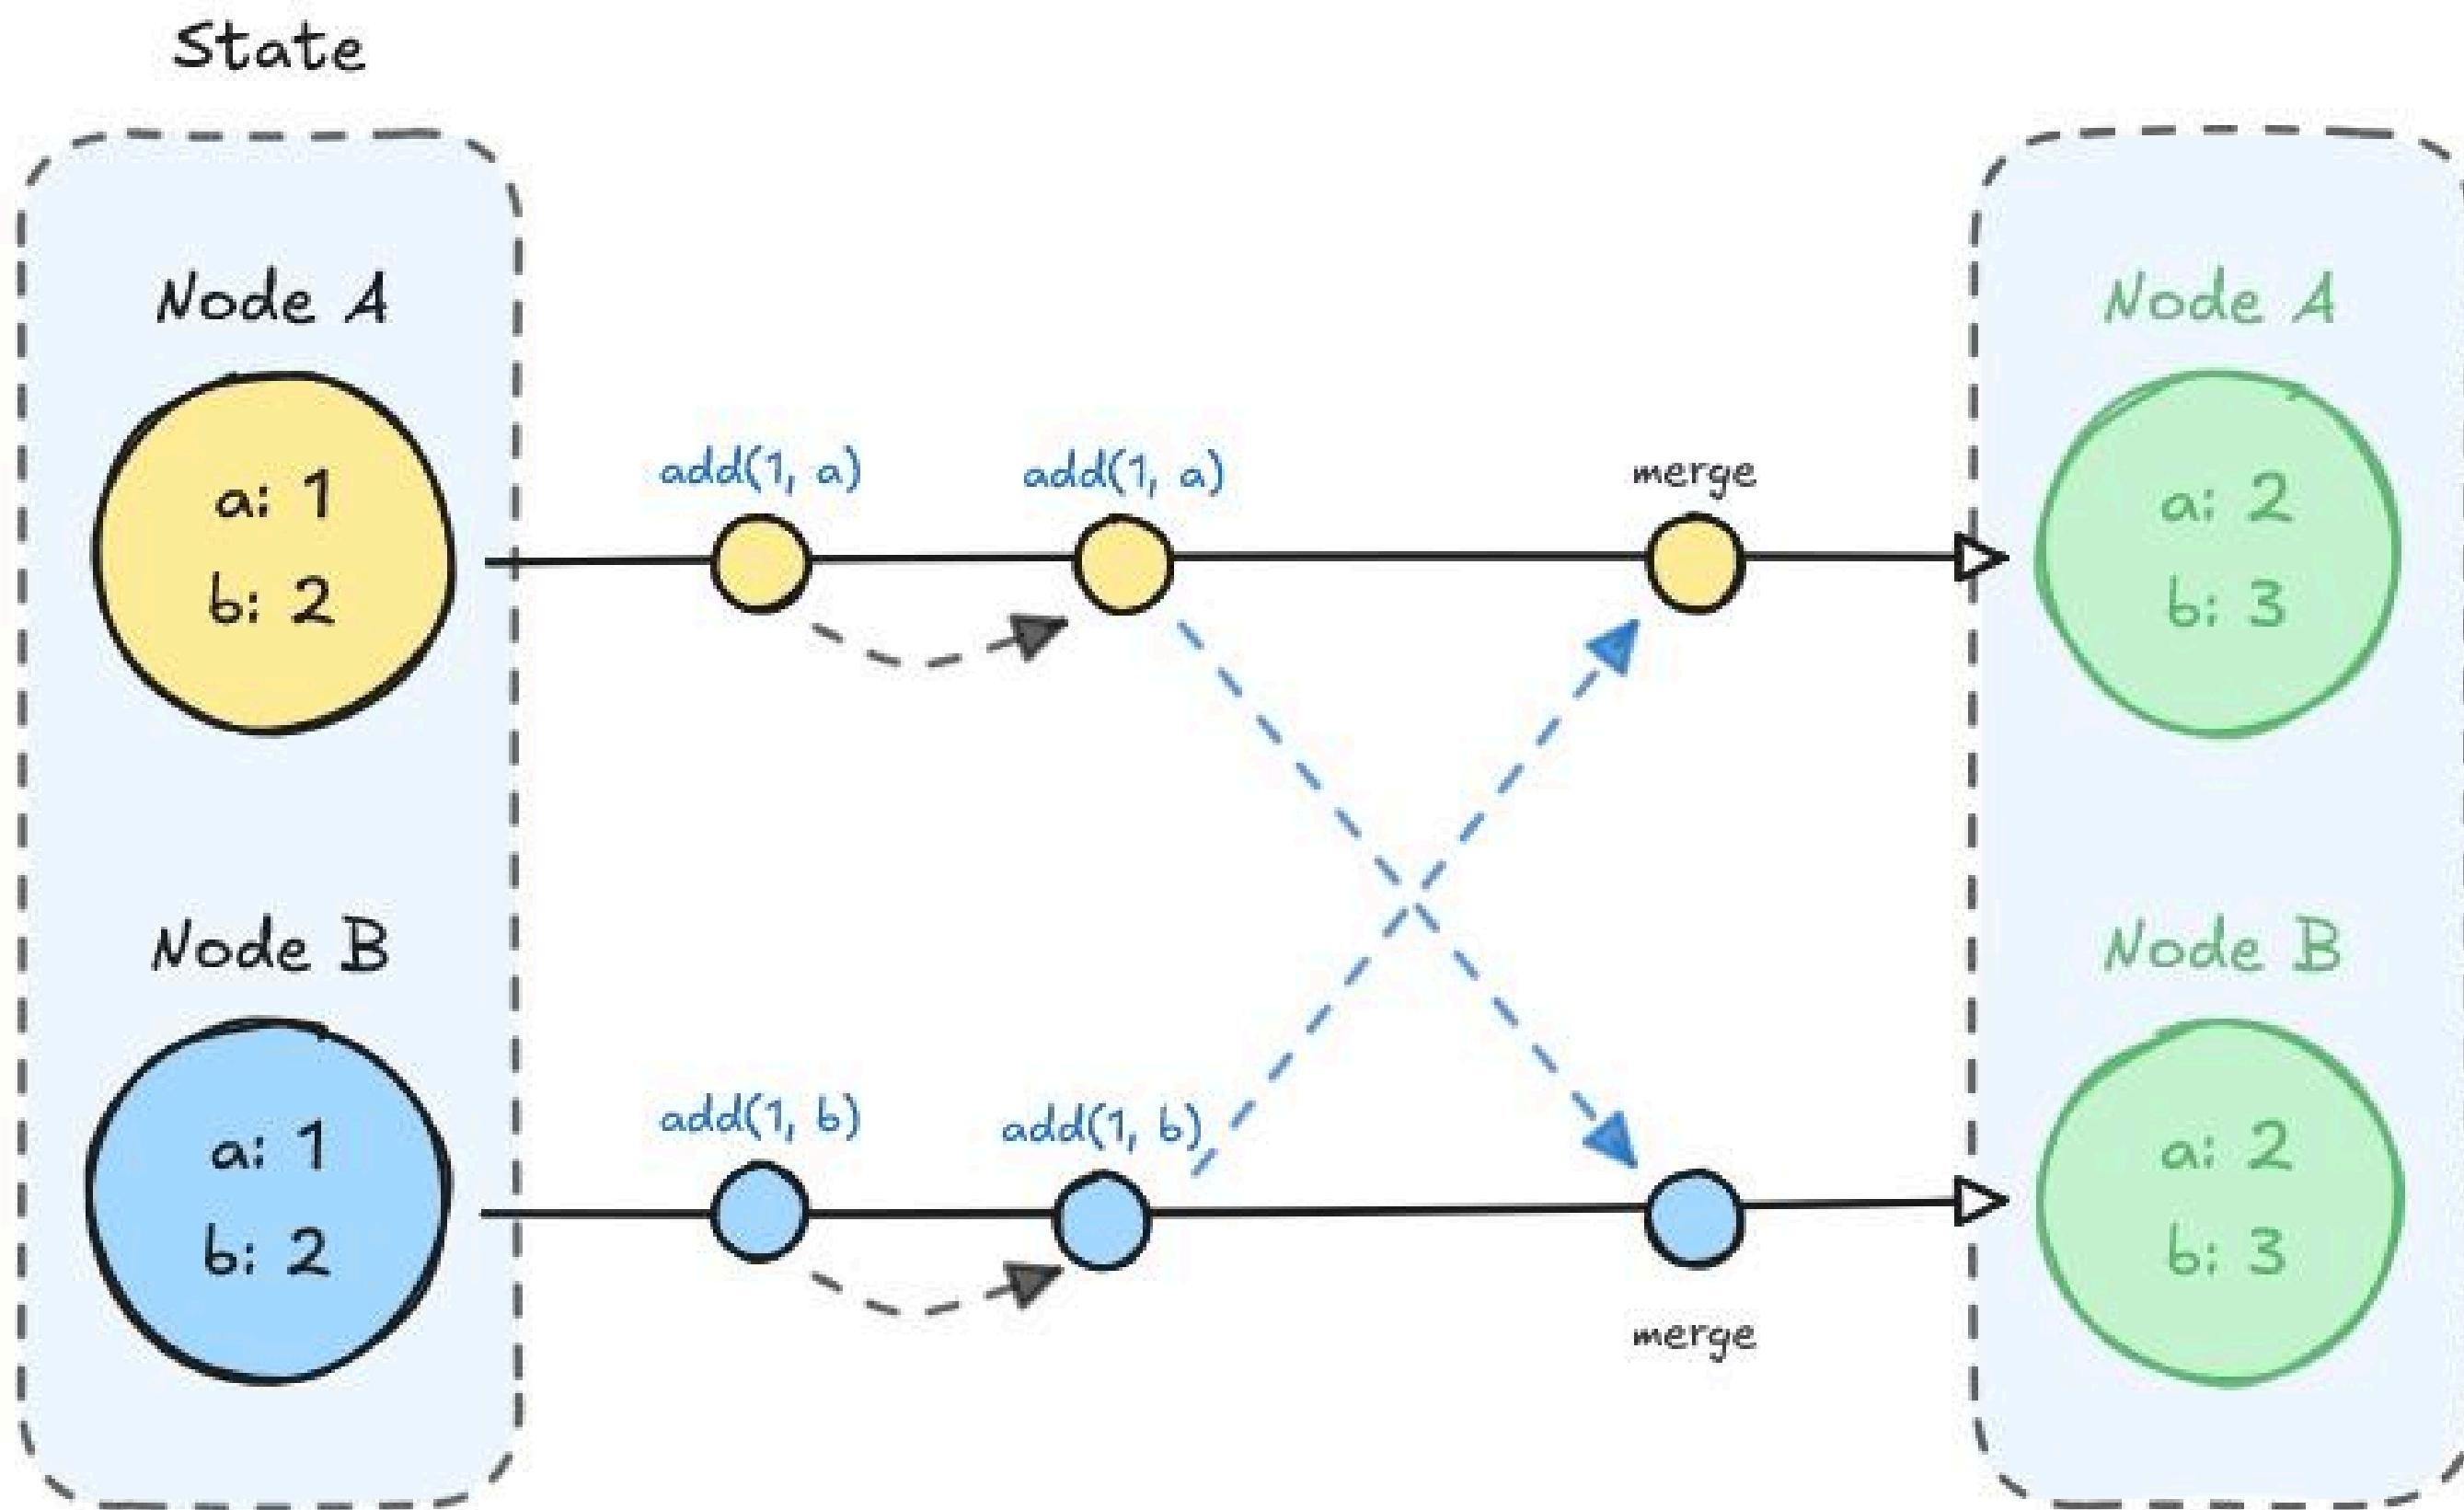
Operation-based CRDTs

Broadcast individual operations to other nodes.

δ state-based CRDTs

Transmit a synthetic state containing only the incremental changes.

Types of CRDT



State-based CRDTs

Broadcast the entire state to other nodes.

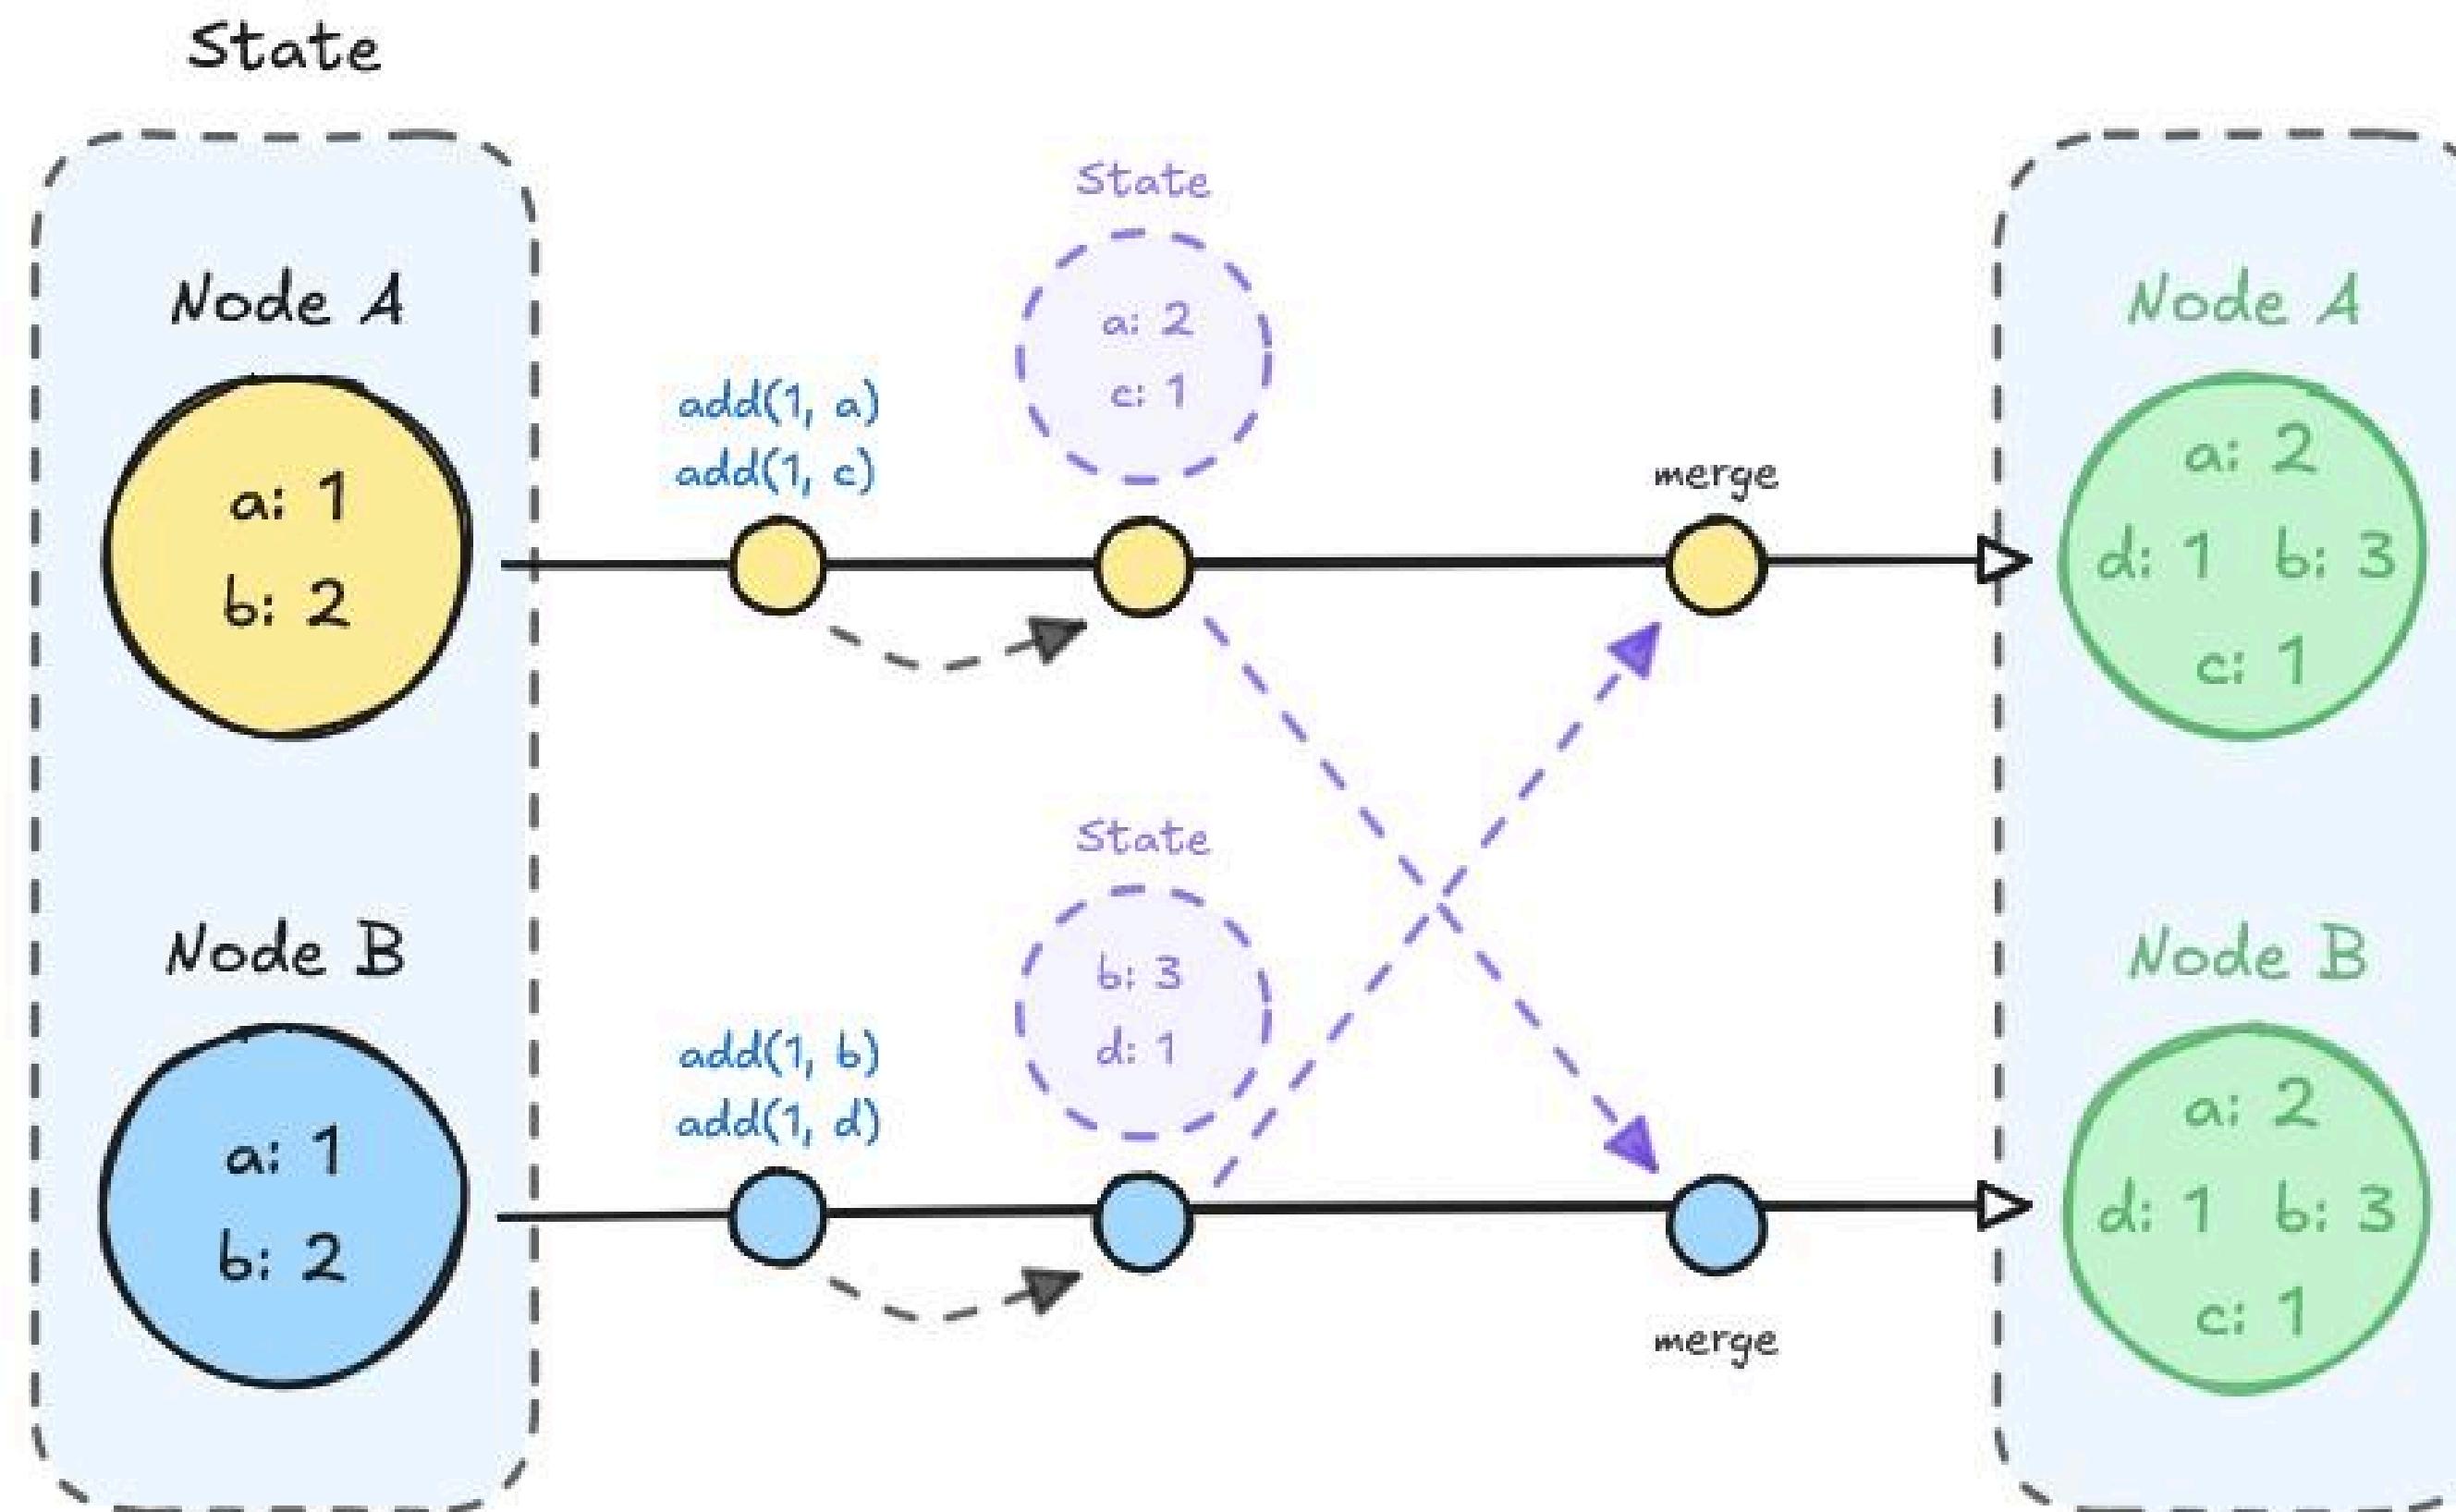
Operation-based CRDTs

Broadcast individual operations to other nodes.

δ state-based CRDTs

Transmit a synthetic state containing only the incremental changes.

Types of CRDT



State-based CRDTs

Broadcast the entire state to other nodes.

Operation-based CRDTs

Broadcast individual operations to other nodes.

δ state-based CRDTs

Transmit a synthetic state containing only the incremental changes.

Advanced CRDT Frameworks



Yjs is a high-performance CRDT framework for building collaborative applications that sync automatically[3].

- It exposes its internal CRDT model as **shared data types** that can be manipulated concurrently.
- It allows for modularity by supporting **several communication protocols** (e.g., *WebRTC*, *WebSockets*, *XMPP*).

[3] P. Nicolaescu, K. Jahns, M. Derntl, and R. Klamma. "Yjs: A Framework for Near Real-Time P2P Shared Editing on Arbitrary Data Types". In *15th International Conference on Web Engineering*. ICWE 2015. Springer LNCS volume 9114, June 2015, pp. 675–678. doi: 10.1007/978-3-319-19890-3_55.

Building a Replicated File Storage System using CRDTs



Key features

- **File insertion and removal**
- On-demand **digital signature validation**
- **Peer-to-peer synchronization**
- **File retention policy** (to prevent storage overhead)

Building a Replicated File Storage System using CRDTs

CRDT Storage

Map-like data type that holds the raw files and their metadata (including their digital signature)

Backend

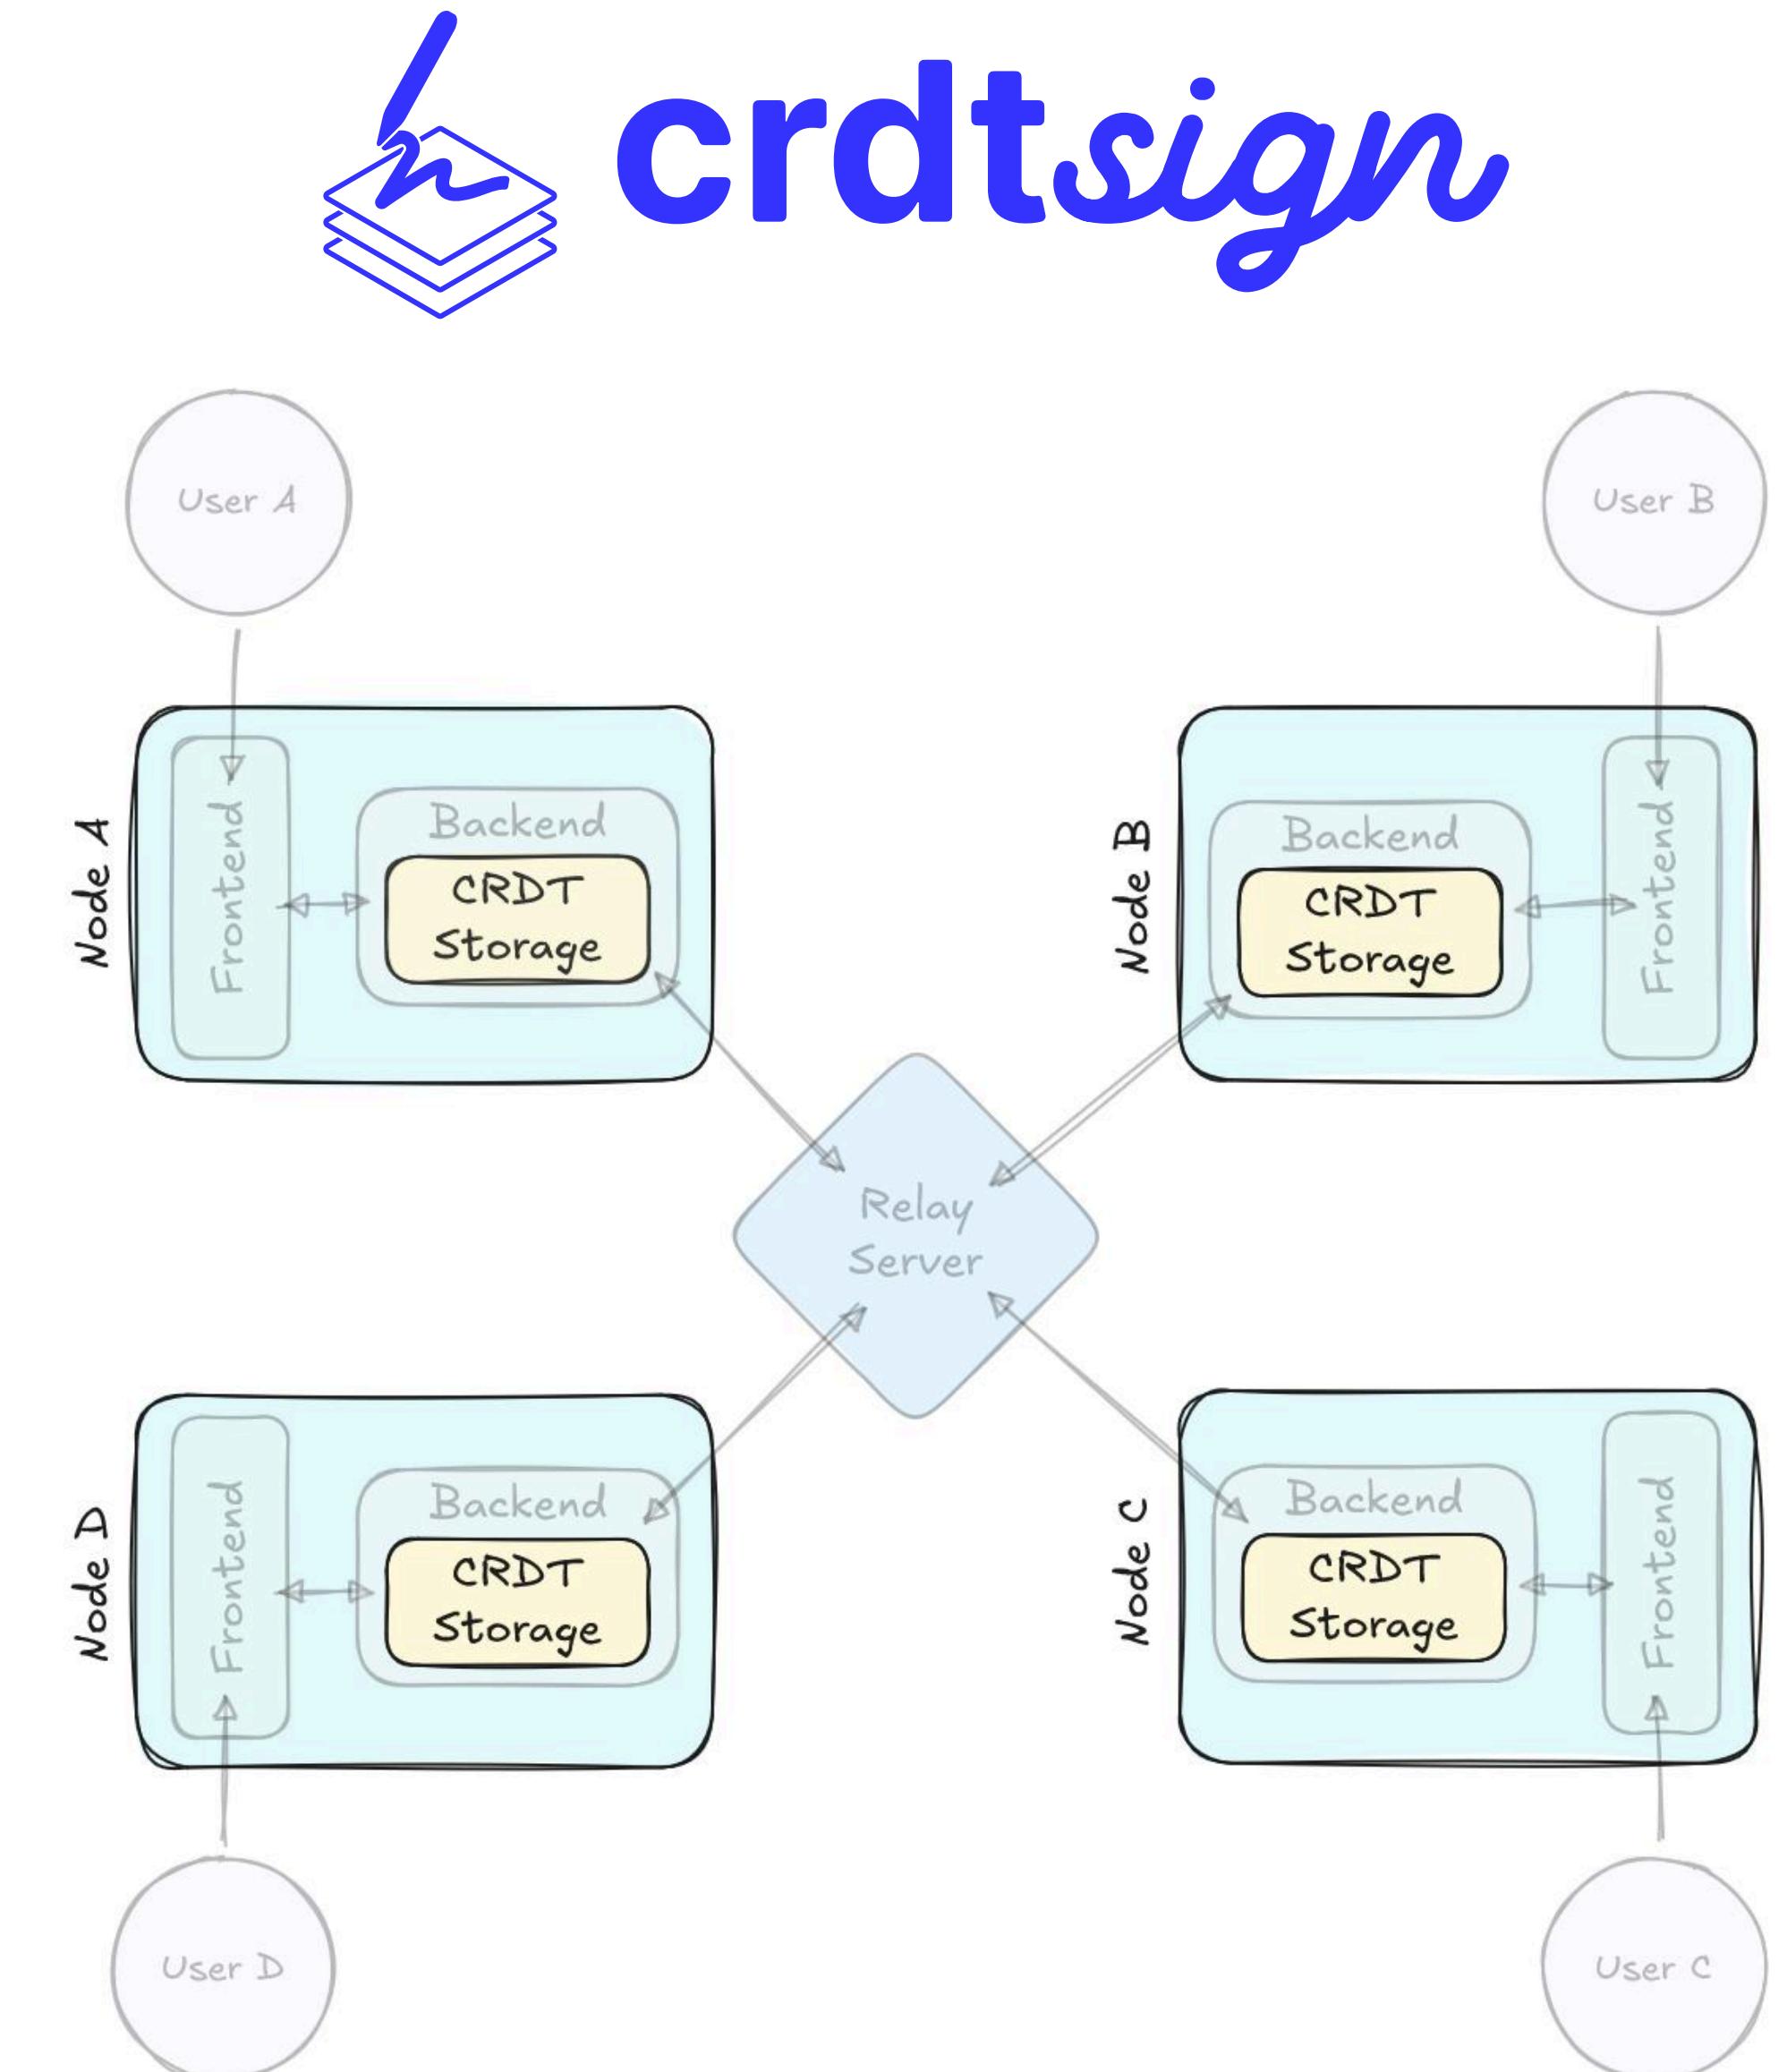
Manages the *CRDT Storage* and the *File Signature* module

Relay Server

Simulates a P2P environment by conveying updates from one node to the others

Frontend

Provided to the *User* as a way to interact with the *Backend*



Building a Replicated File Storage System using CRDTs

CRDT Storage

Map-like data type that holds the raw files and their metadata (including their digital signature)

Backend

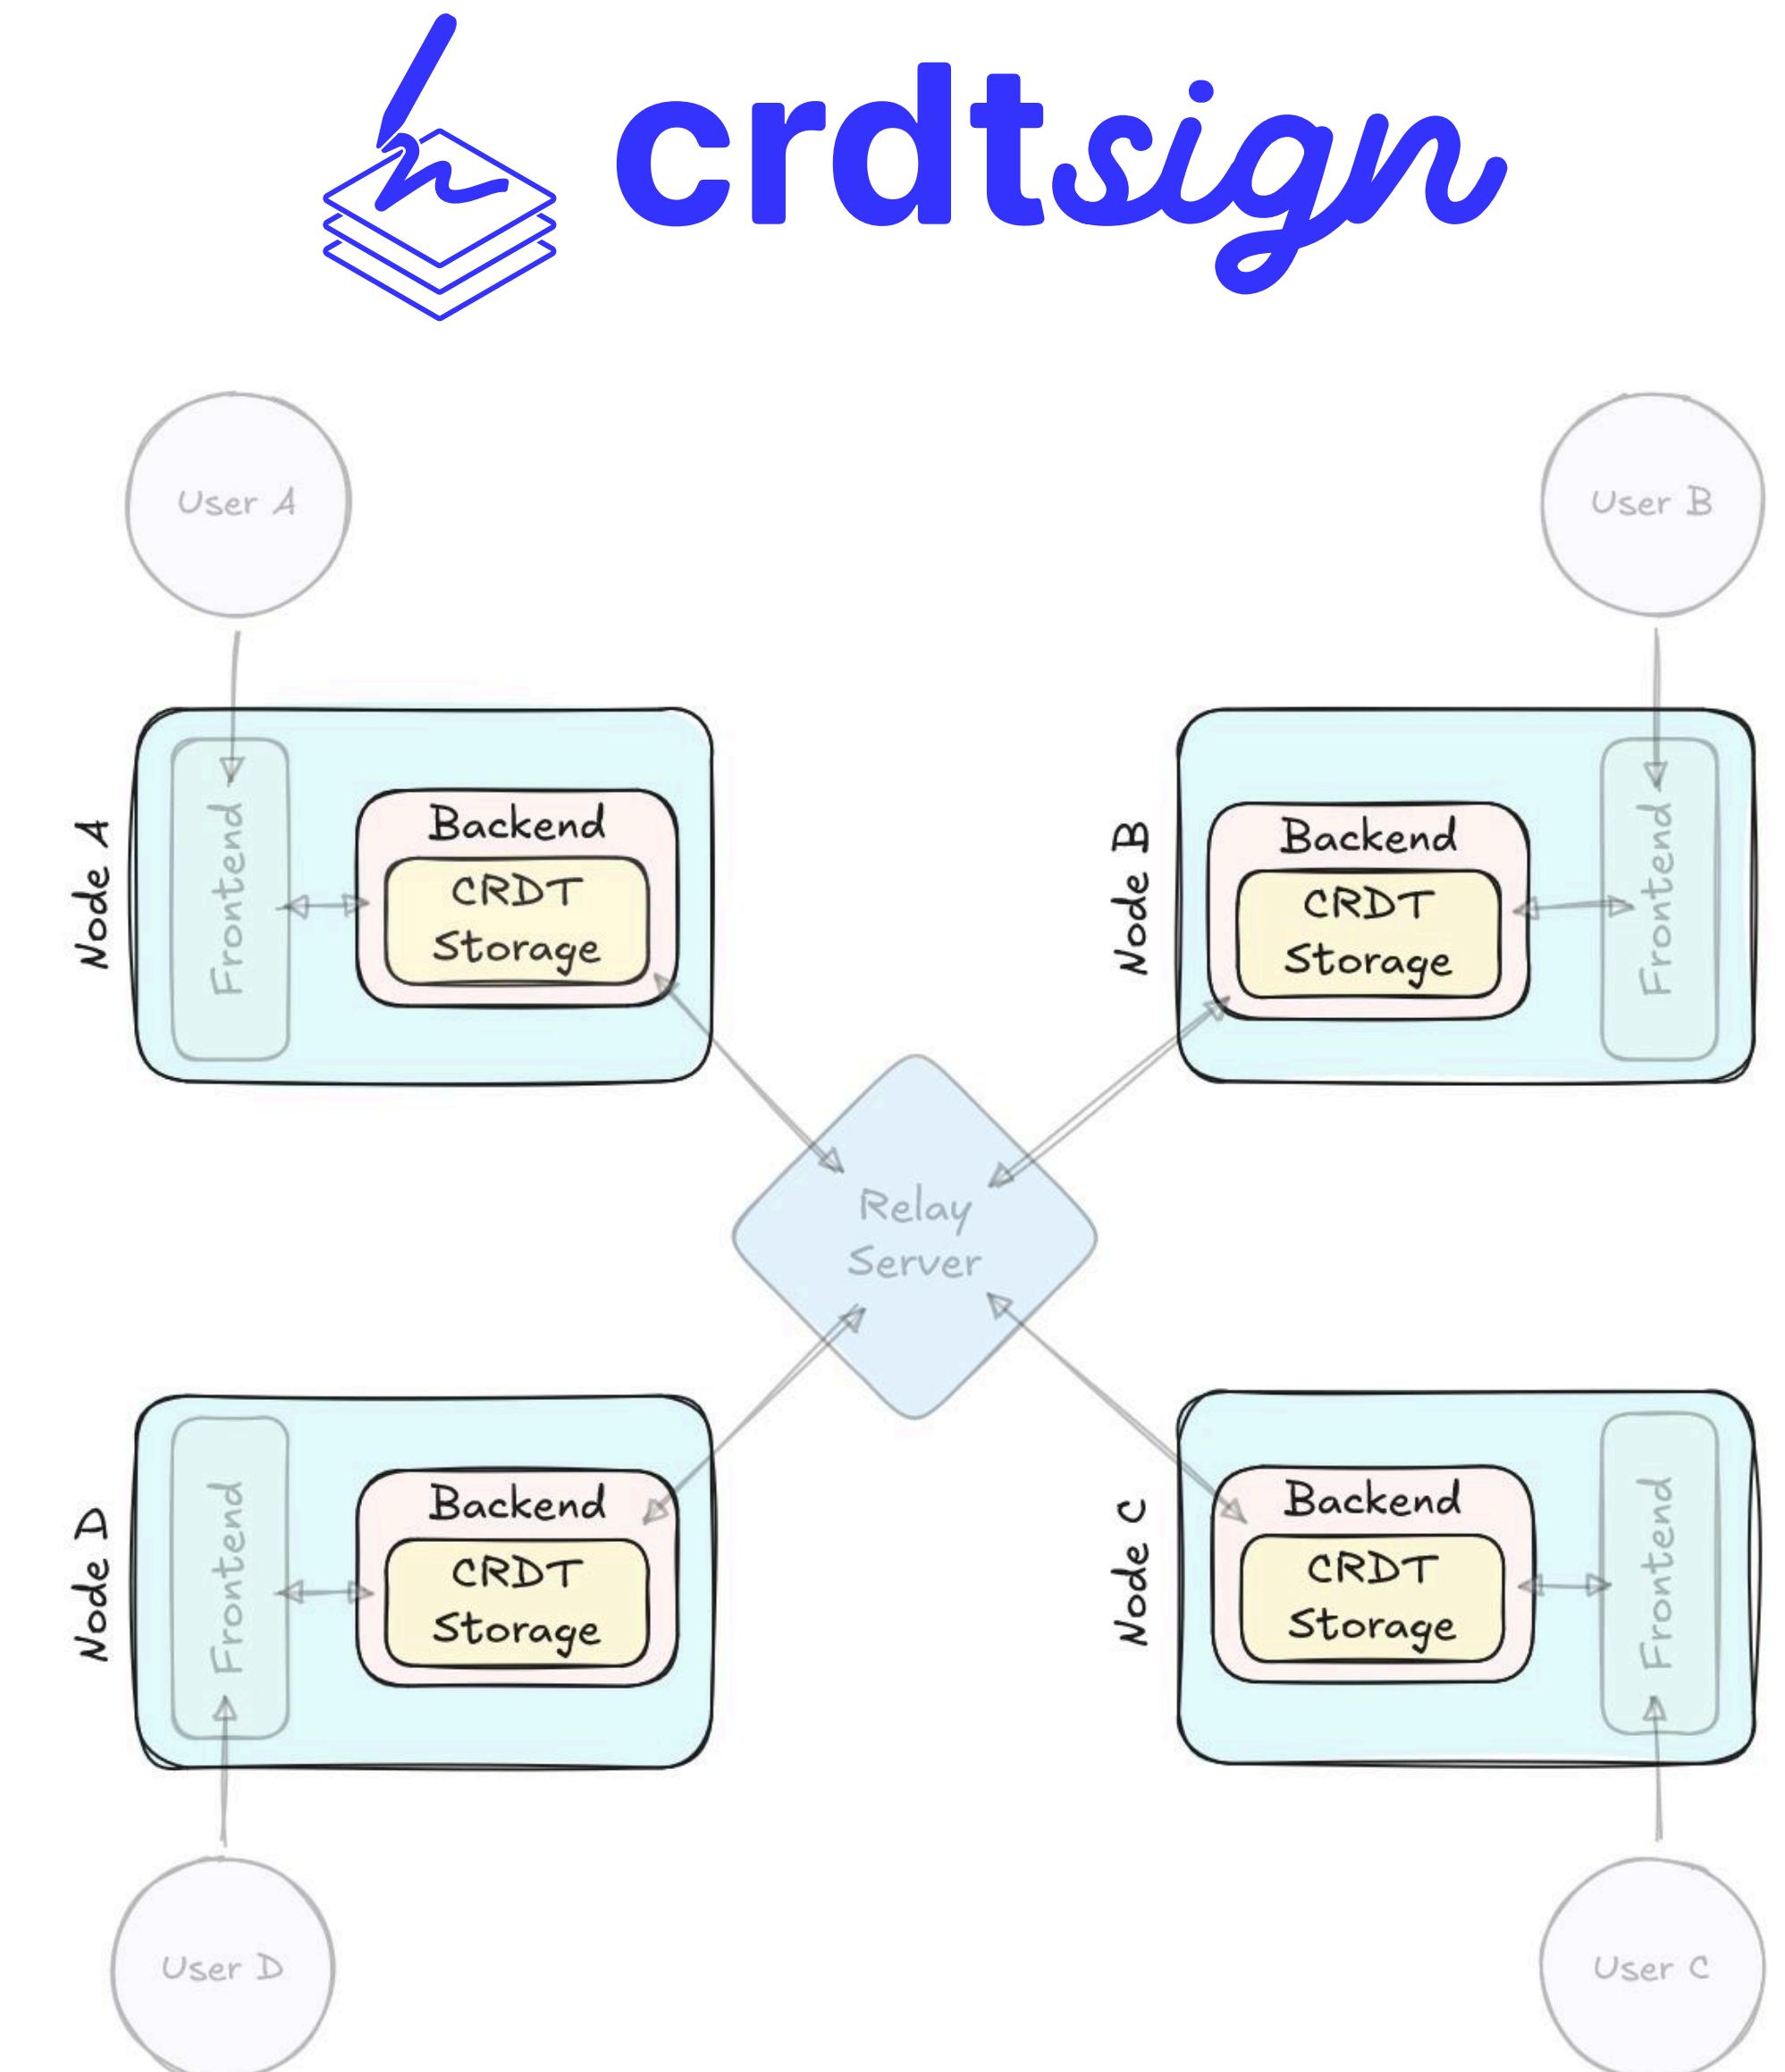
Manages the *CRDT Storage* and the *File Signature* module

Relay Server

Simulates a P2P environment by conveying updates from one node to the others

Frontend

Provided to the *User* as a way to interact with the *Backend*



Building a Replicated File Storage System using CRDTs

CRDT Storage

Map-like data type that holds the raw files and their metadata (including their digital signature)

Backend

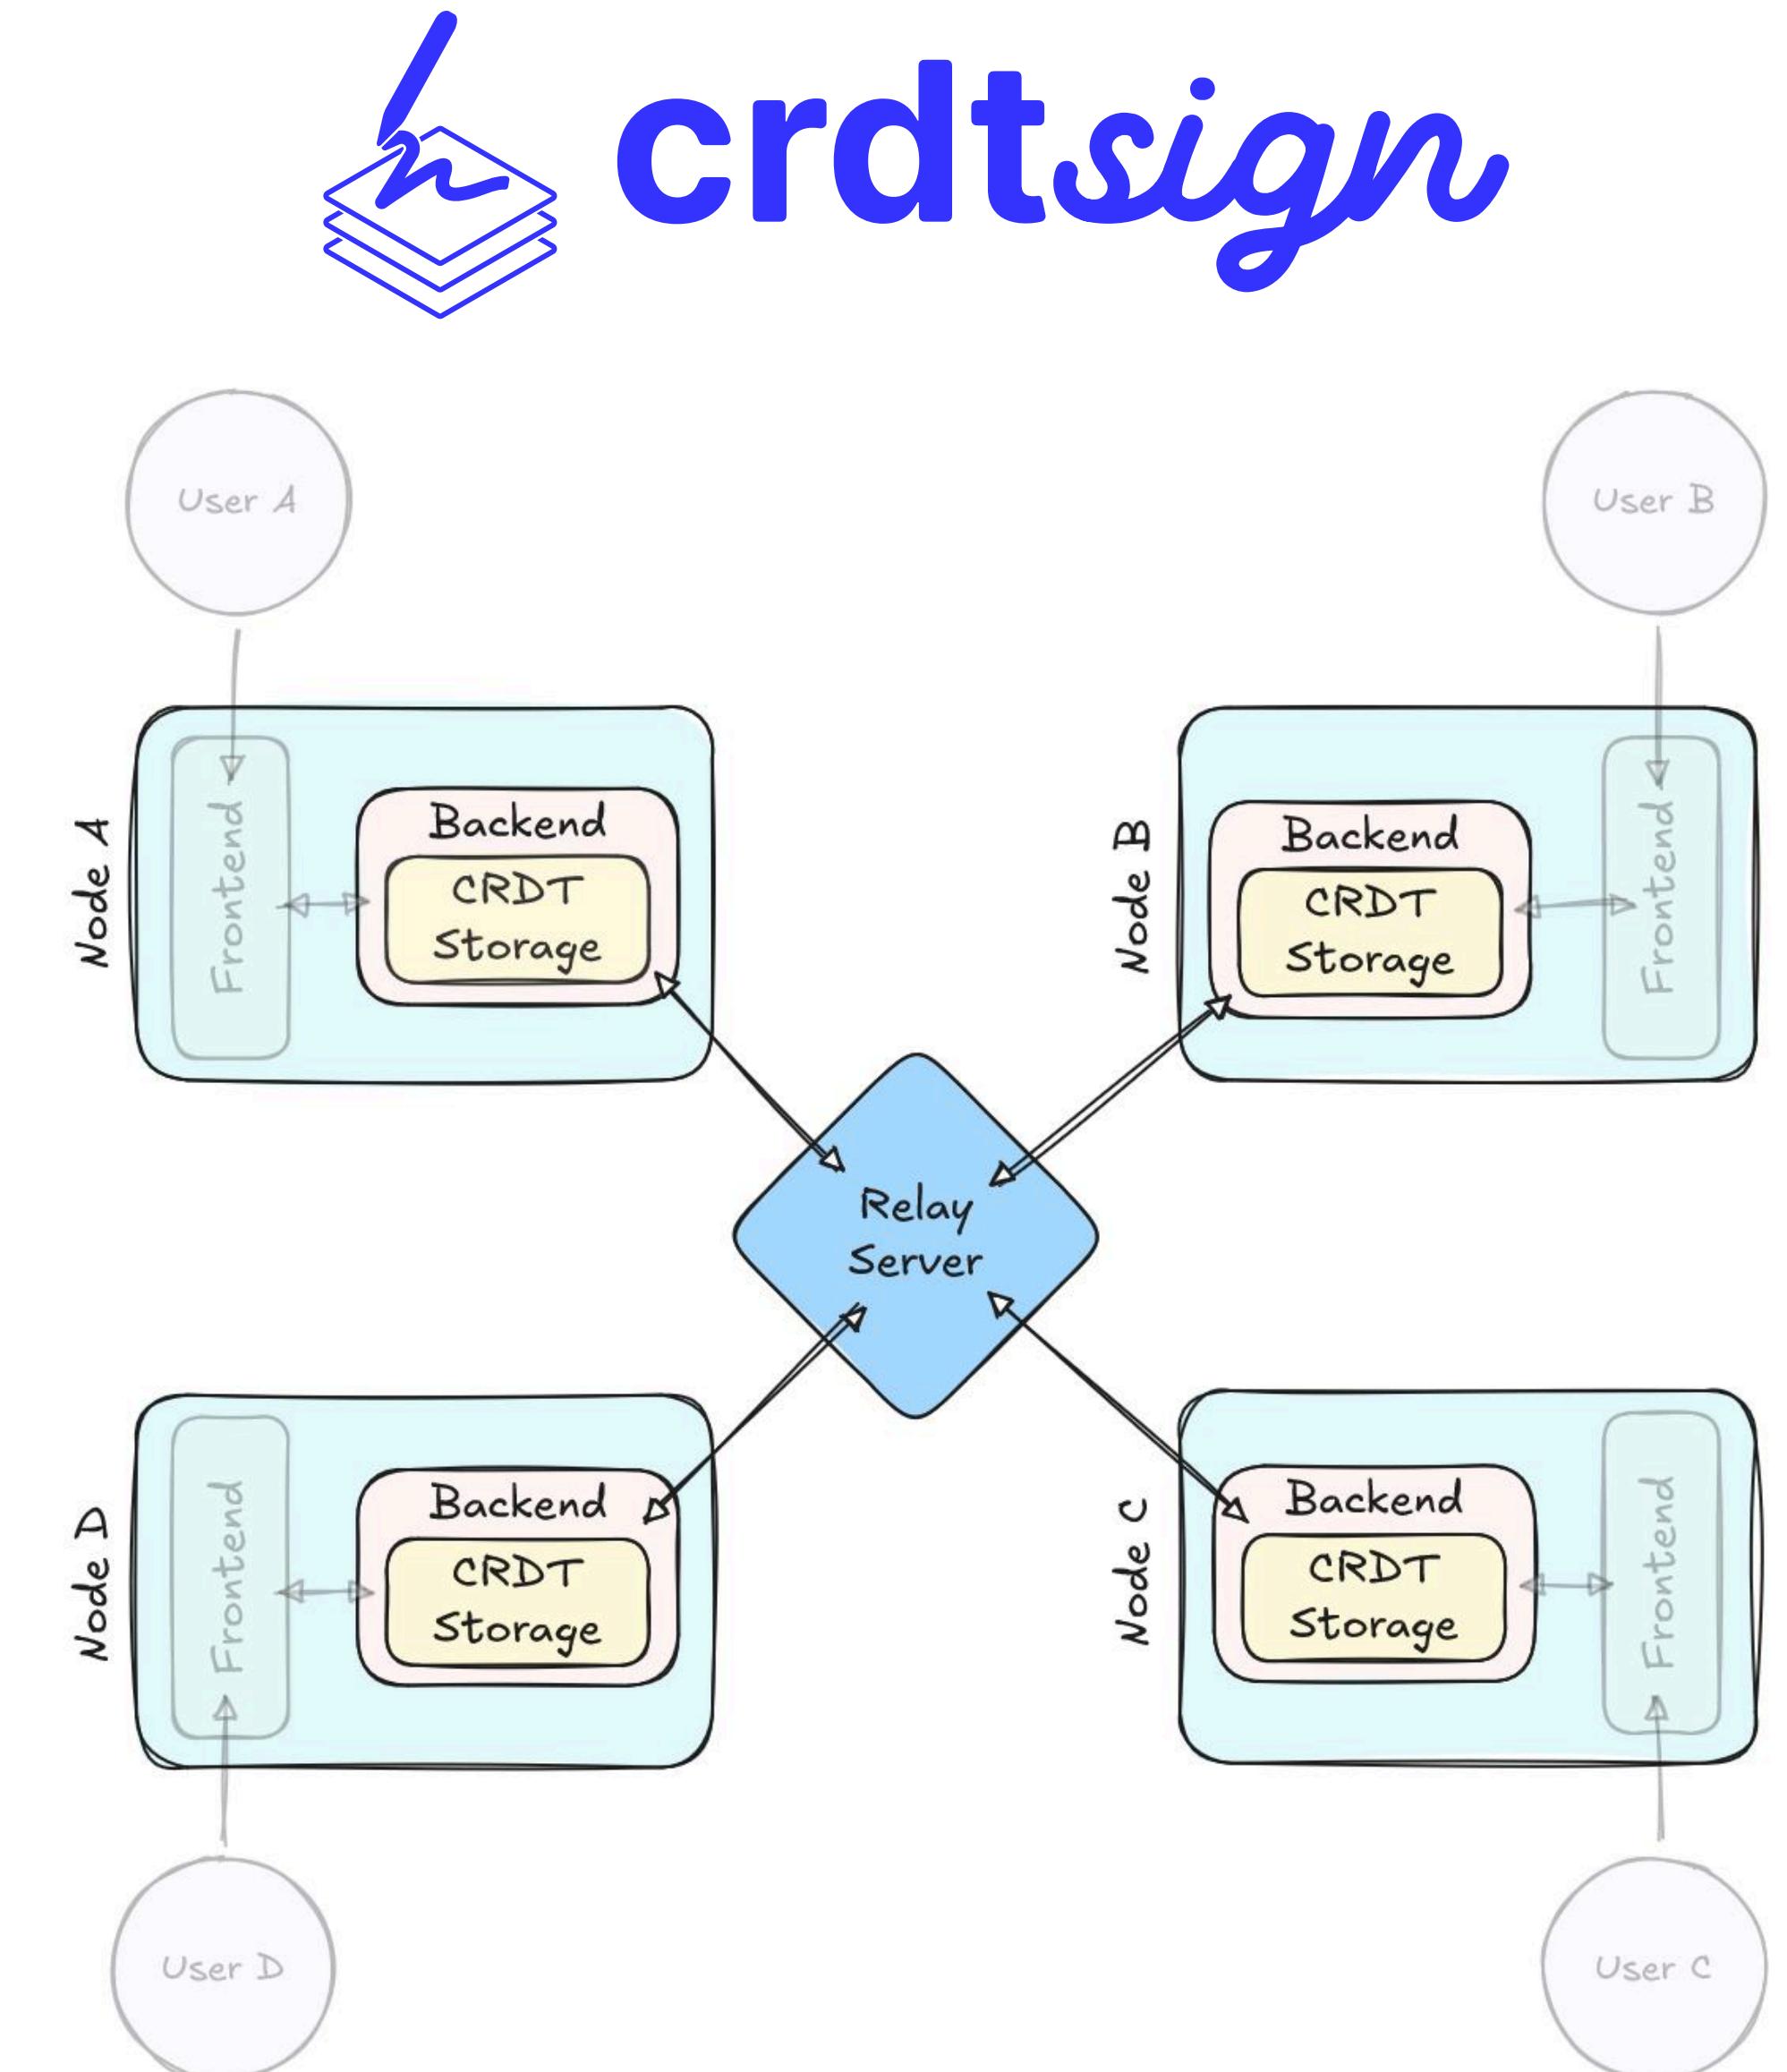
Manages the *CRDT Storage* and the *File Signature* module

Relay Server

Simulates a P2P environment by conveying updates from one node to the others

Frontend

Provided to the *User* as a way to interact with the *Backend*



Building a Replicated File Storage System using CRDTs

CRDT Storage

Map-like data type that holds the raw files and their metadata (including their digital signature)

Backend

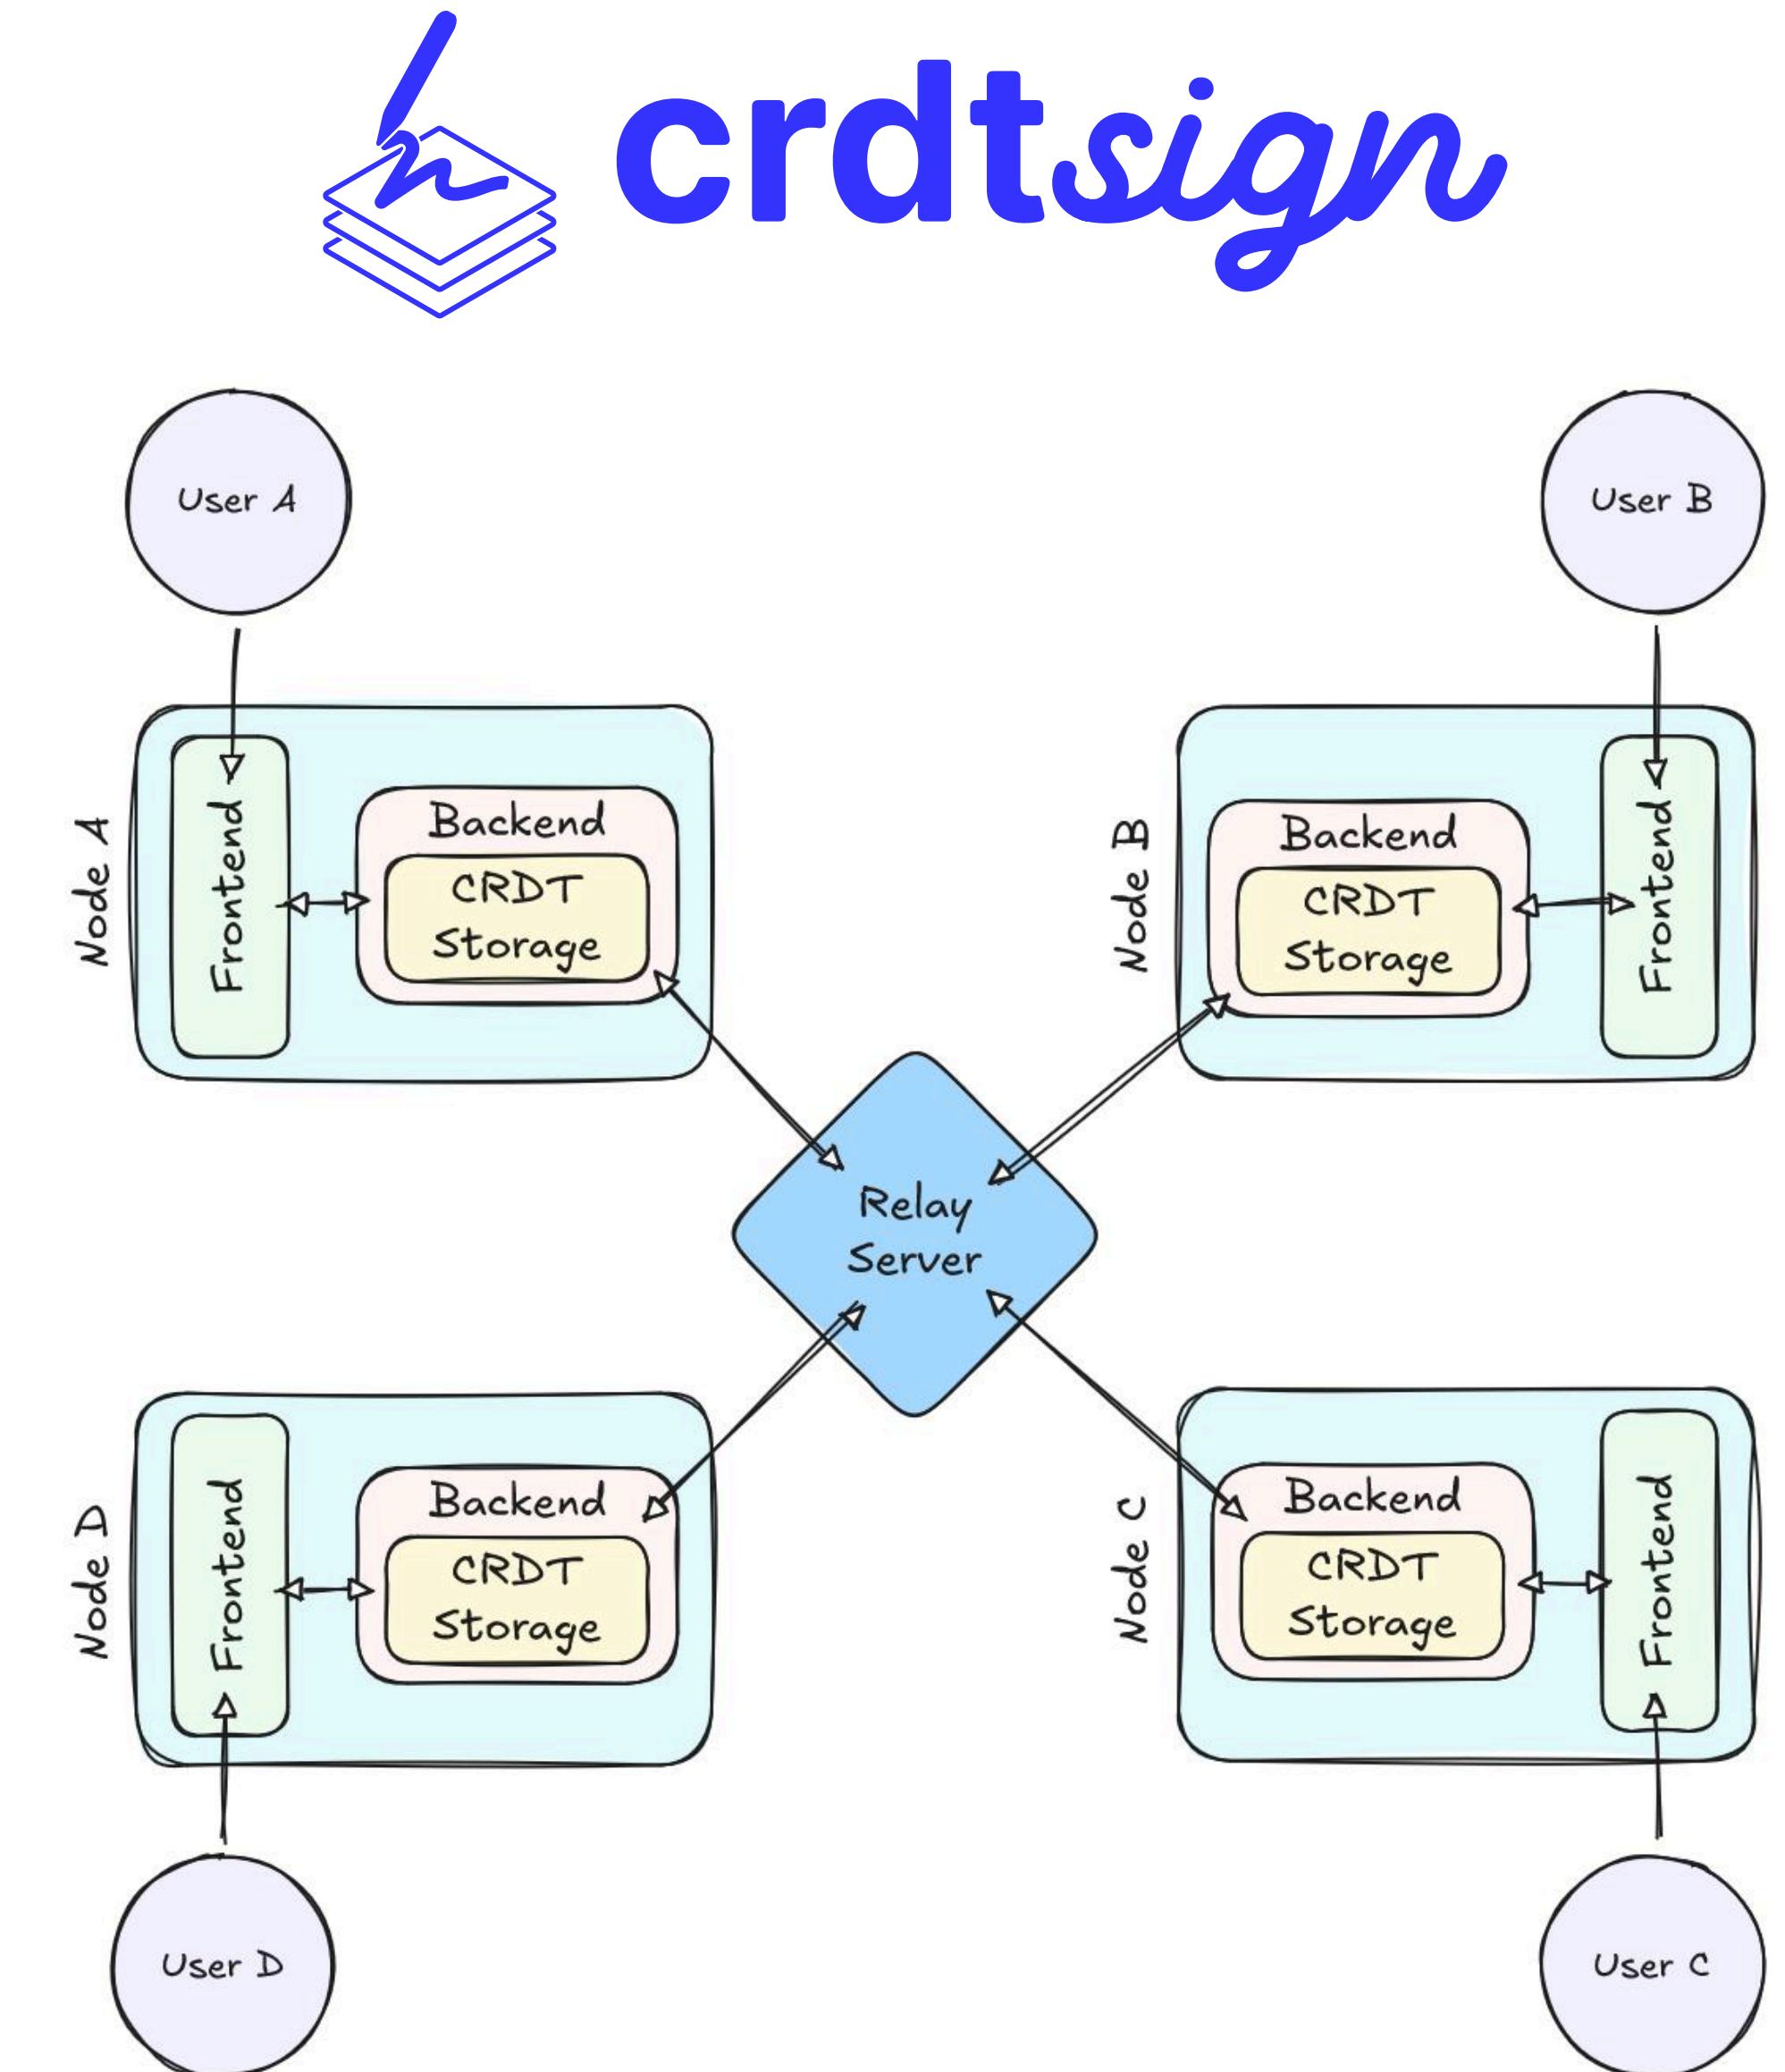
Manages the *CRDT Storage* and the *File Signature* module

Relay Server

Simulates a P2P environment by conveying updates from one node to the others

Frontend

Provided to the *User* as a way to interact with the *Backend*

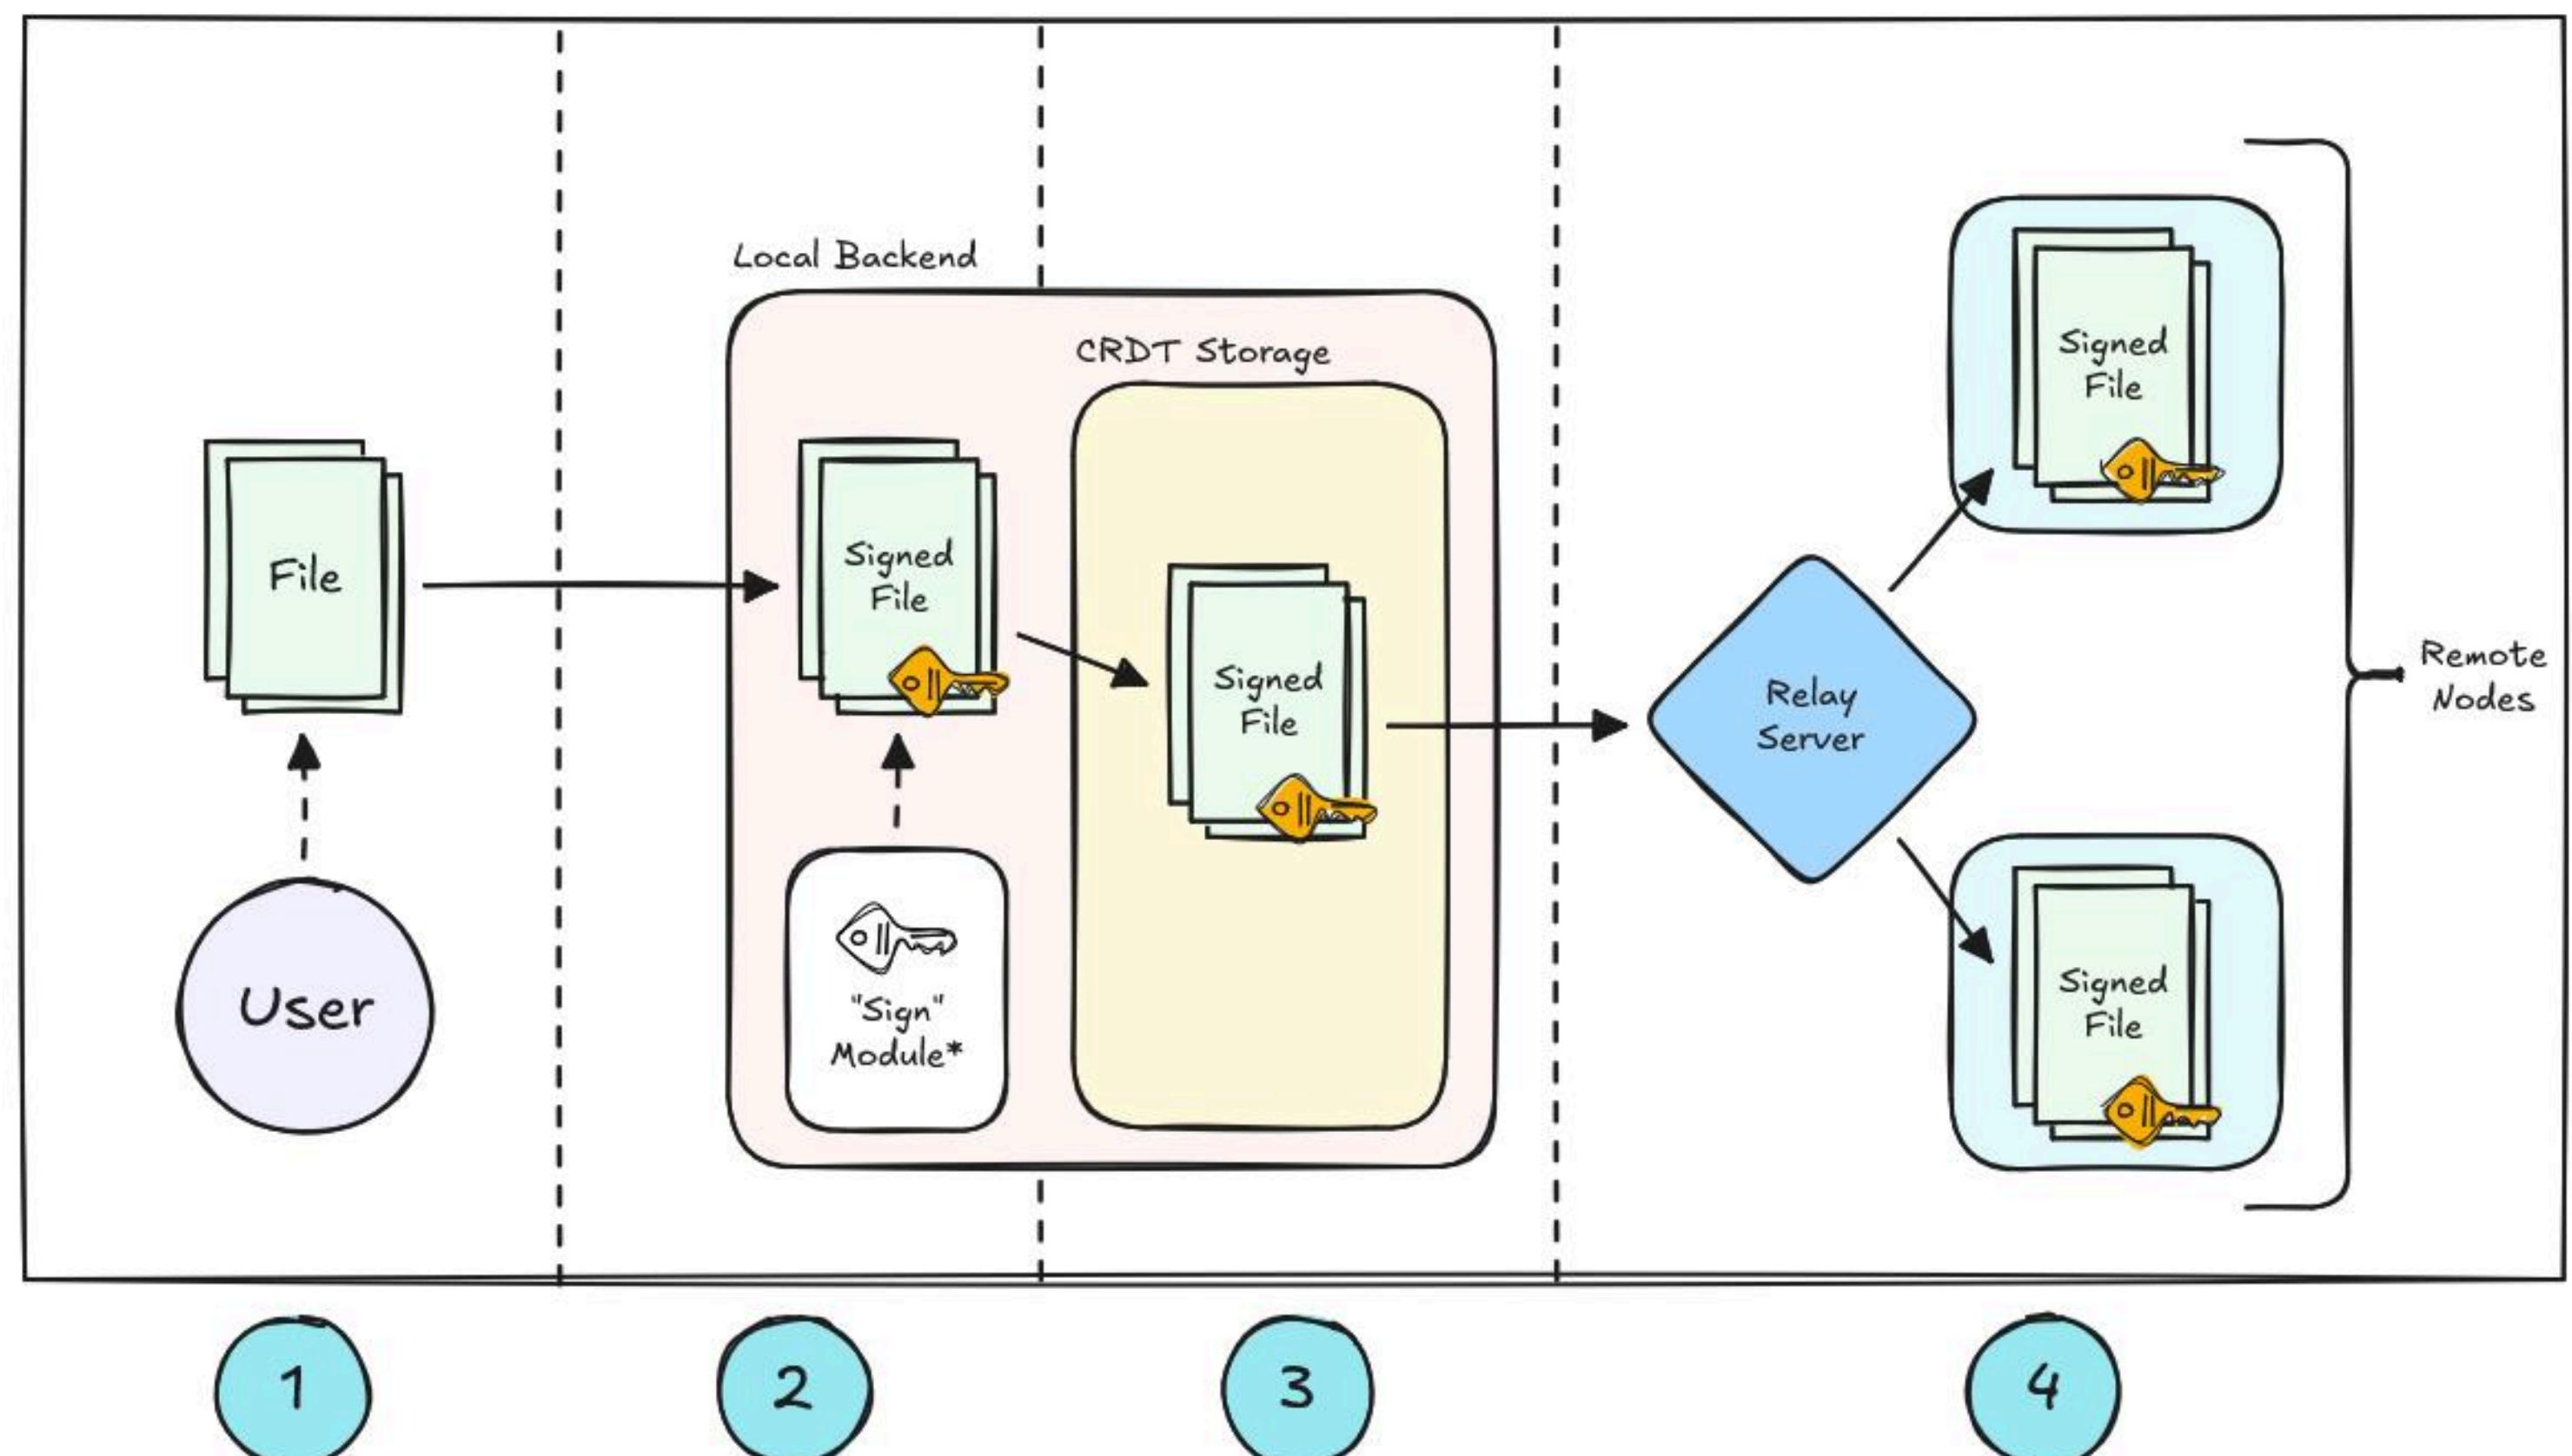


EXAMPLE - FILE INSERTION

Building a Replicated File Storage System using CRDTs

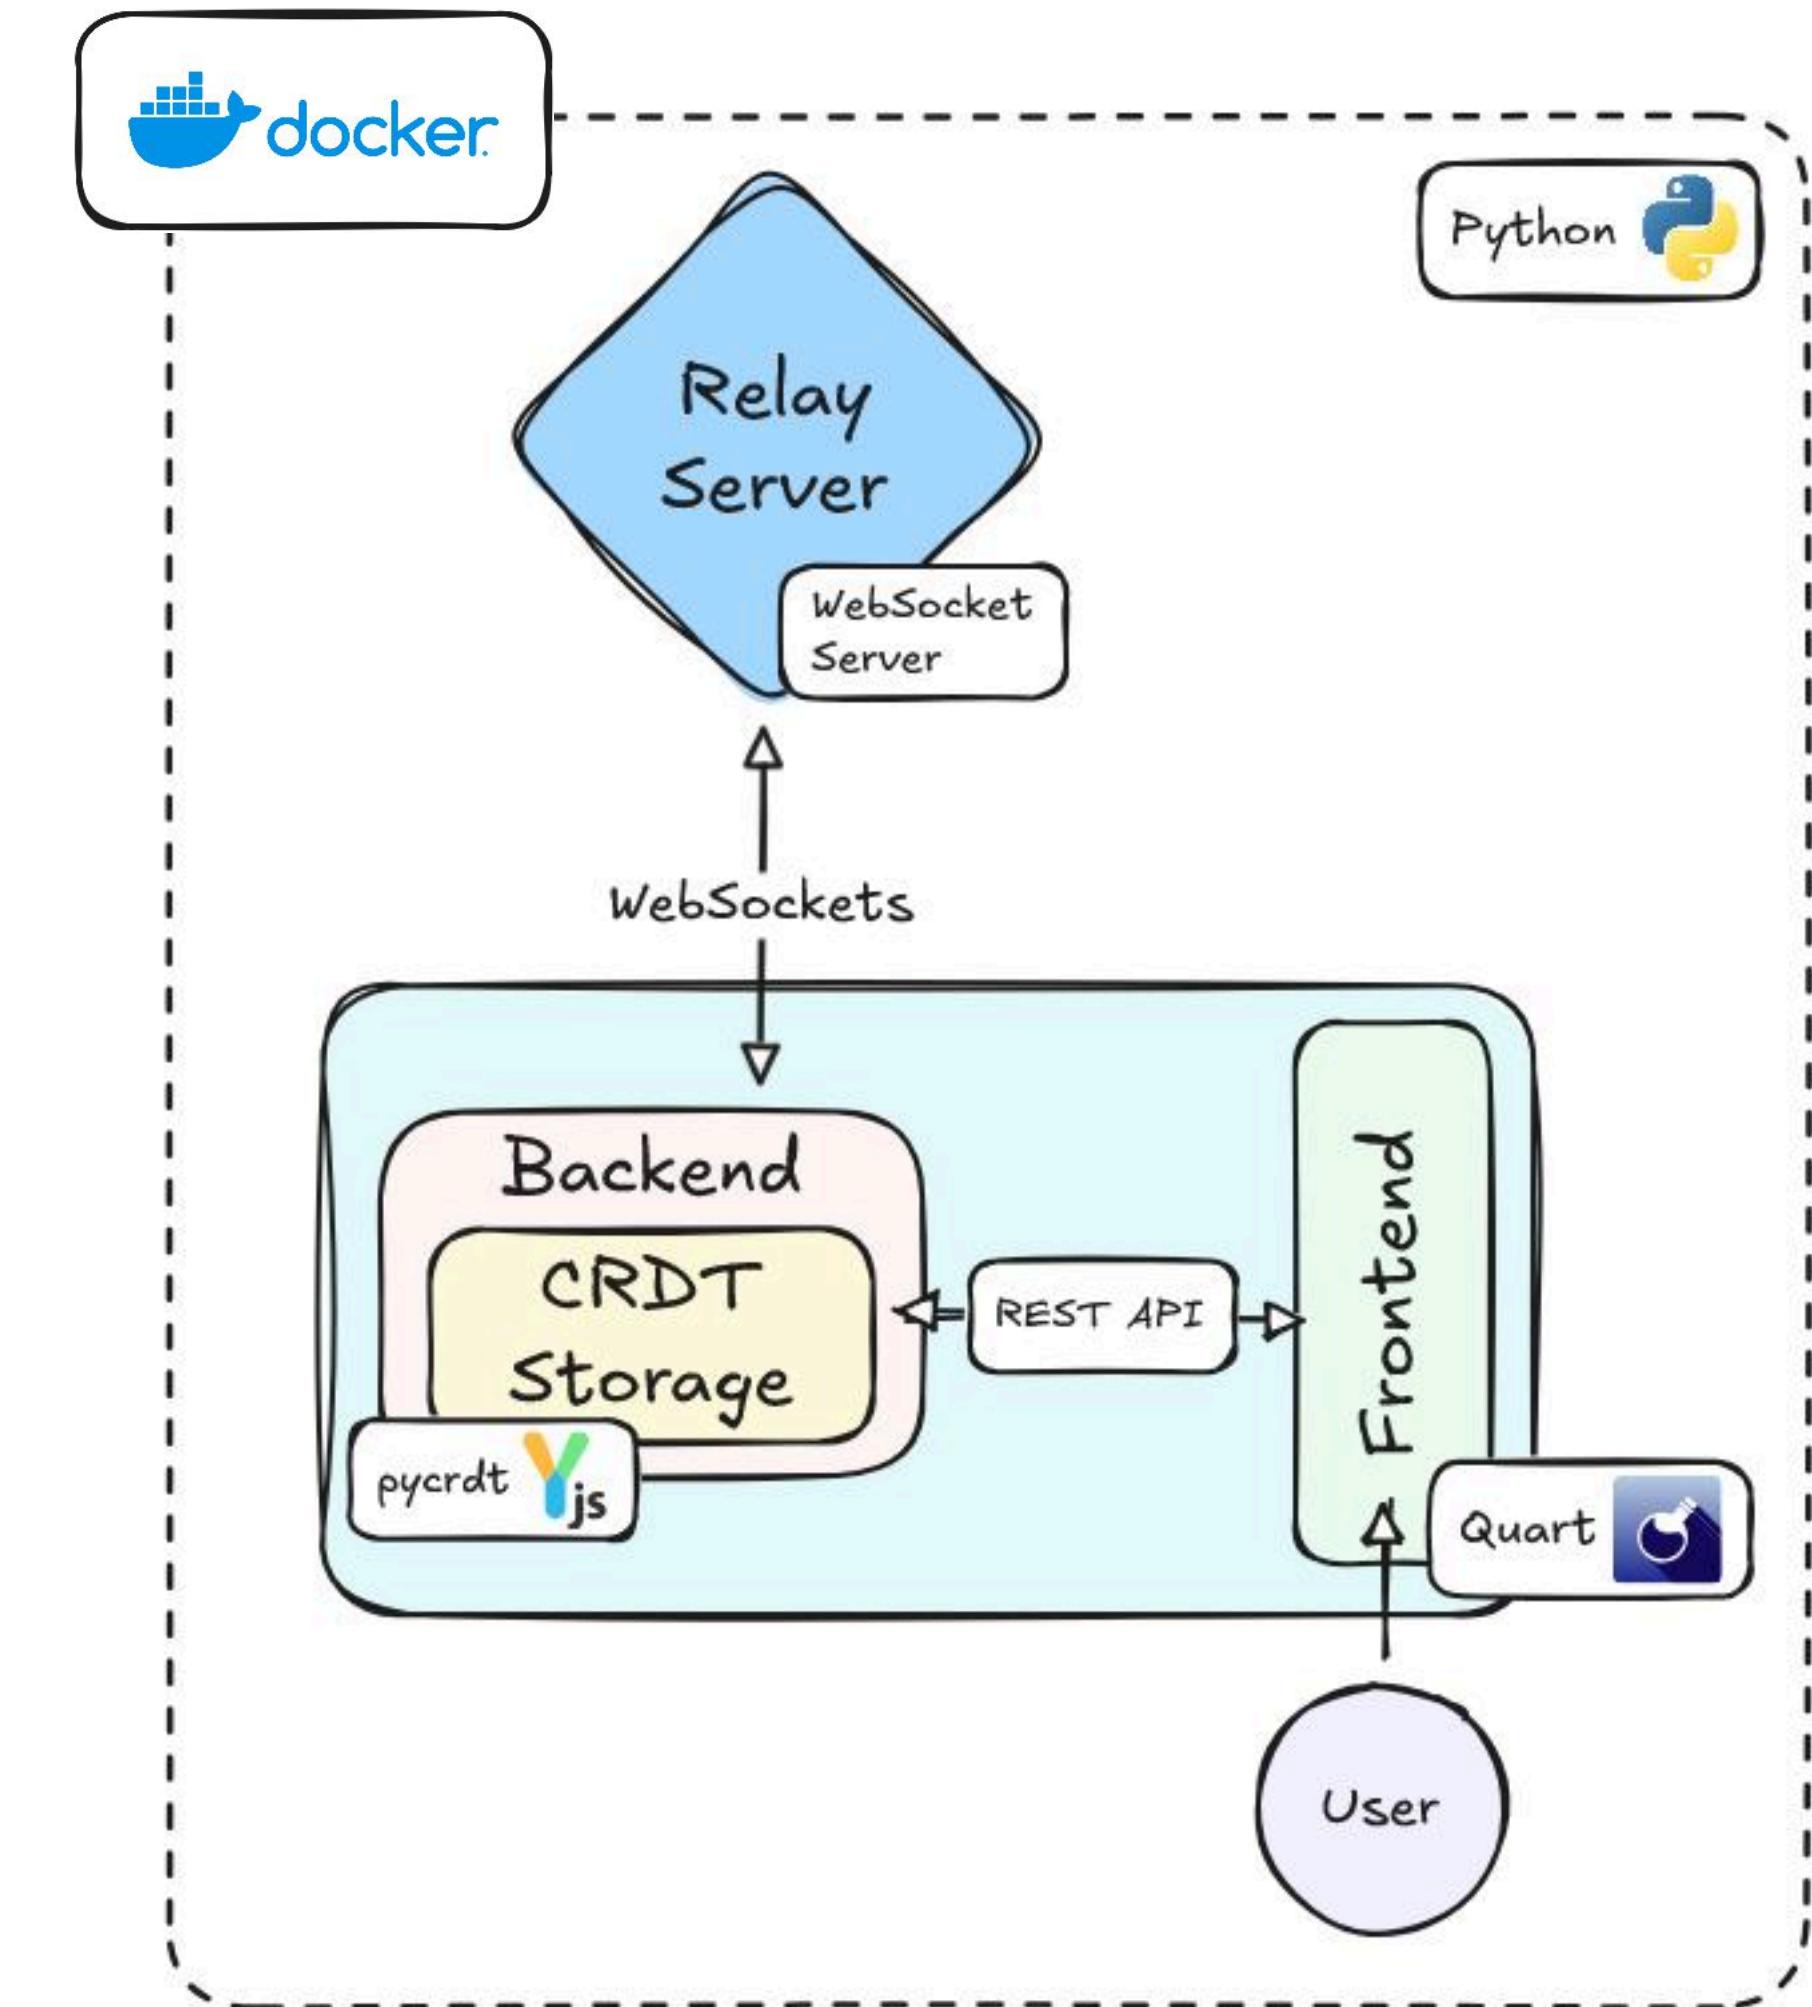


1. The user **uploads** a file
2. The Backend **sends** the file to other nodes (via the Relay Server)
3. The Backend **saves** the file to CRDT storage
4. The Backend **sends** the file to other nodes (via the Relay Server)



Building a Replicated File Storage System using CRDTs

- Written in **Python**
- **Backend** (library): *pycrdt* library (*Yjs* for Python)
- **Frontend**: *REST API* accessed client-side through a *Web application* (written in *HTML* and *JavaScript*, deployed with *Quart*)
- Communication with **Relay Server**: *WebSockets* protocol
- **Deployment**: Docker (each component has its own container)



WEB APPLICATION

Building a Replicated File Storage System using CRDTs



The screenshot shows the 'Signed Files' section of the crdt sign web application. It displays a table with three rows of signed files:

FILE NAME	USERNAME	SIGNED ON	EXPIRATION DATE	ACTIONS
document-1.pdf	User A	1/21/2026, 9:45:29 AM	1/21/2026, 9:55:00 AM	Details Validate Delete
document-2.pdf	User B	1/21/2026, 9:46:50 AM	1/21/2026, 9:56:00 AM	Details Validate Delete
document-3.pdf	User C	1/21/2026, 9:47:03 AM	1/21/2026, 9:56:00 AM	Details Validate Delete

The screenshot shows the 'Upload a File' section of the crdt sign web application. The form fields are as follows:

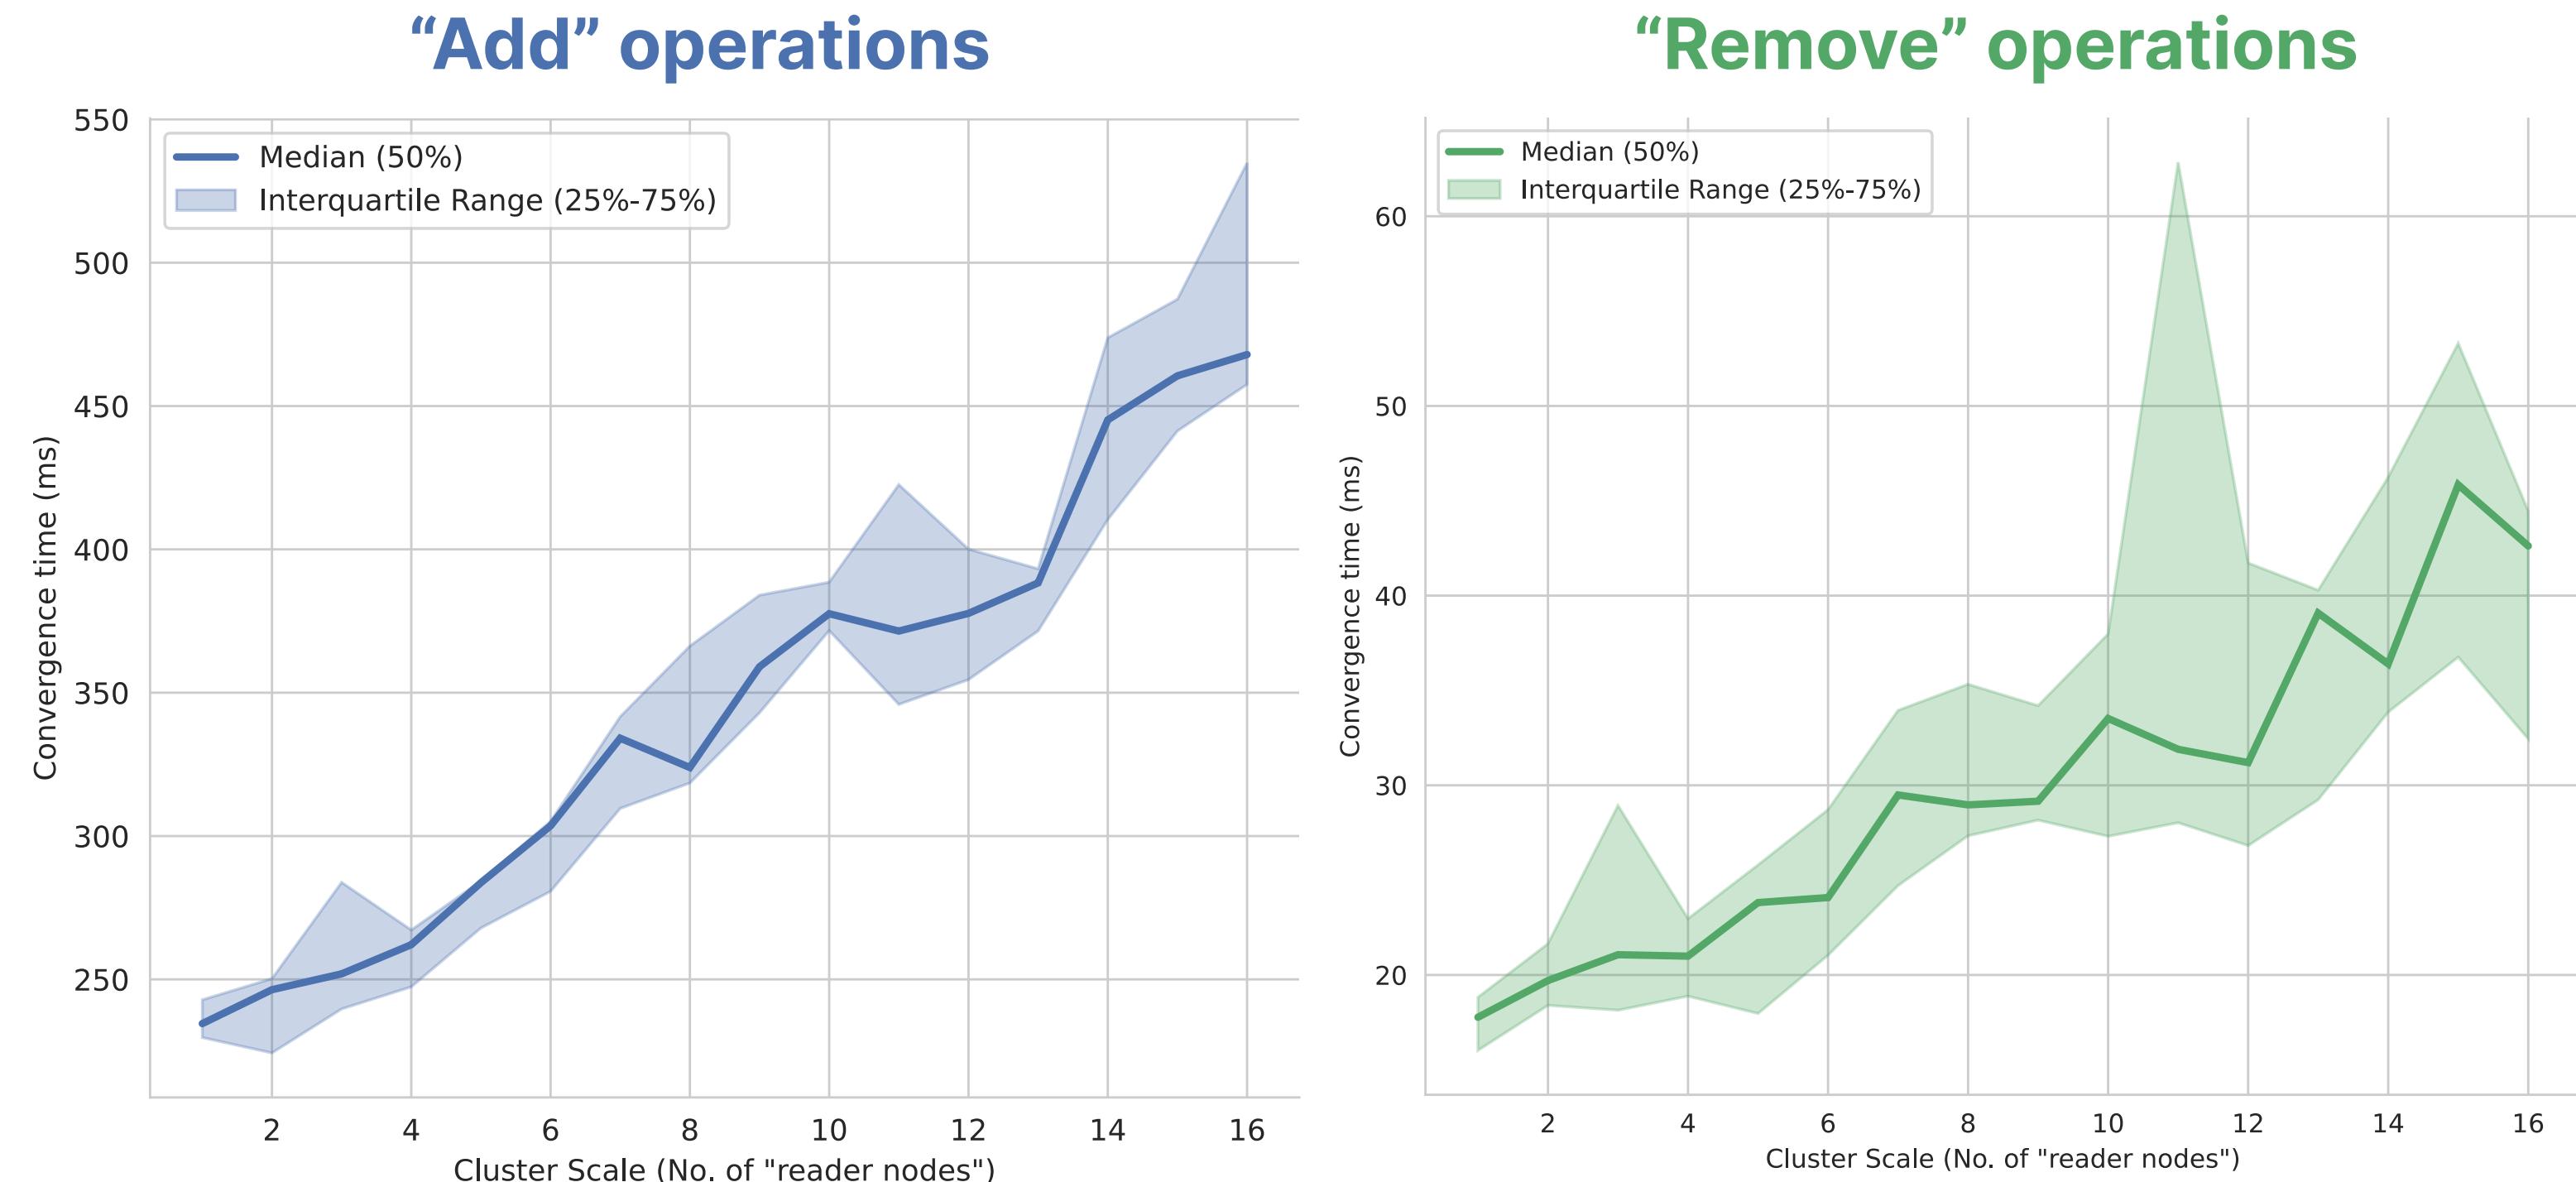
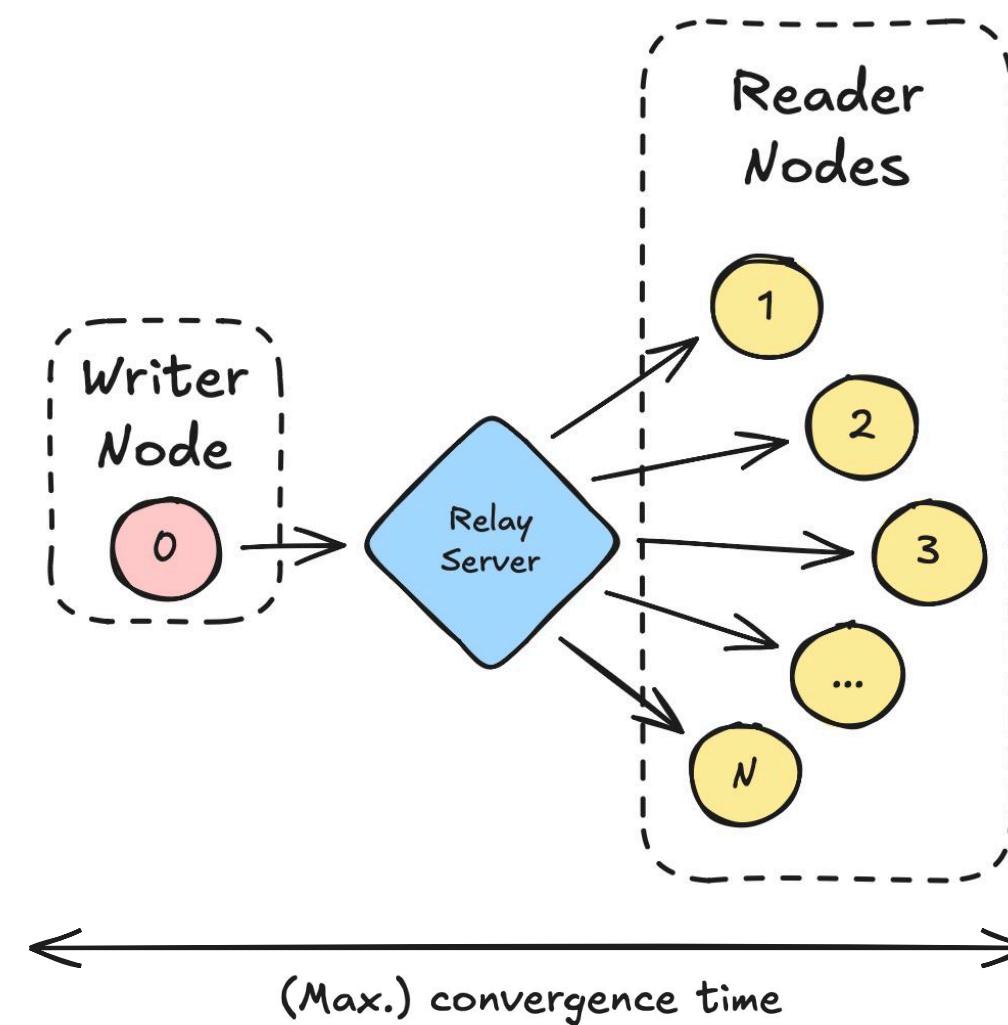
- Select File: Choose File document-3.pdf
- User ID: user_z6KTJYPhvdmr
- Username: User C
- Expiration Date: 21/01/2026, 09:56

Below the form, a green success message states: "File signed successfully! Signature: 641e412c6c66c61f5f65..."

Building a Replicated File Storage System using CRDTs



Experimental Setup



Overall, the system maintains <1 second (1000ms) response times across all tested scales, meeting user expectations for responsiveness.

Conclusion

Contributions of the Thesis

- **Replicated File Storage**
 - Basic functionalities (*file insertion and removal*)
 - Blockchain-like traceability

Limitations of the proof-of-concept

- File embedded within the CRDT map (suboptimal)
- Test performance limited by experiment hardware

Commercial CRDT-based products

- *Redis*
- Distributed databases (*Riak, AntidoteDB*)

Possibility for future improvements

- Implement “true” P2P protocols (e.g., *Dat*)
- Extended file storage functionalities (e.g., folders, file renaming, ...)
- Large-scale experimental setup (over multiple geographic locations)

Thank you for your attention!